



Philosophical Review

Solving the Skeptical Problem

Author(s): Keith DeRose

Source: *The Philosophical Review*, Vol. 104, No. 1 (Jan., 1995), pp. 1-52

Published by: Duke University Press on behalf of Philosophical Review

Stable URL: <http://www.jstor.org/stable/2186011>

Accessed: 07-01-2018 14:51 UTC

JSTOR is a not-for-profit service that helps scholars, researchers, and students discover, use, and build upon a wide range of content in a trusted digital archive. We use information technology and tools to increase productivity and facilitate new forms of scholarship. For more information about JSTOR, please contact support@jstor.org.

Your use of the JSTOR archive indicates your acceptance of the Terms & Conditions of Use, available at <http://about.jstor.org/terms>



Philosophical Review, *Duke University Press* are collaborating with JSTOR to digitize, preserve and extend access to *The Philosophical Review*

Solving the Skeptical Problem

Keith DeRose

1. The Puzzle of Skeptical Hypotheses

Many of the most celebrated, intriguing, and powerful skeptical arguments proceed by means of skeptical hypotheses. Brutally pared to their barest essentials, they are roughly of the following form, where 'O' is a proposition about the external world one would ordinarily think one knows (e.g., I have hands¹) and 'H' is a suitably chosen skeptical hypothesis (e.g., I am a bodiless brain in a vat who has been electrochemically stimulated to have precisely those sensory experiences I've had, henceforth a 'BIV'²):

The Argument from Ignorance (AI)³

1. I don't know that not-H.

2. If I don't know that not-H, then I don't know that O.

So, C. I don't know that O.⁴

¹I choose this O partly for its historical connections to Descartes's First Meditation, and also because I think it is an exemplary case of something we ordinarily think we know. But while we would ordinarily think we know this O, we'd seldom have occasion to say that we know it, because cases in which such a claim to knowledge would be conversationally in order are quite rare. (Exception: A teacher begins an epistemology lecture by matter-of-factly listing various things she knows, and that any plausible theory of knowledge should make her come out to know. In the course of this listing, she says, "And I know that I have hands.") For this and various related reasons, some might not like my choice of O. Such readers are invited to supply their own favorite exemplary cases of things we know as the skeptic's target.

²Those who think that Hilary Putnam may have already disarmed BIV-inspired skepticism should understand the BIV hypothesis to be the hypothesis that one's brain has been *recently* envatted after many years of normal embodiment. For even if Putnam is right in claiming that the content of the beliefs of the BIVs of his scenario is such that these BIVs aren't massively deceived, it seems that recently envatted BIVs are so deceived.

³AI takes its name primarily from its first premise. But since one of AI's best formulations (to which I hereby refer readers seeking a good version of AI that has not been so brutally pared) is in chapter 1 of Peter Unger's book *Ignorance: A Case for Scepticism* (1975), it is in more than one sense that it is an argument "from ignorance."

⁴I actually haven't pared AI to its *barest* essentials. It could be further

Setting aside the distracting side issues that immediately threaten from all directions, and keeping AI in this stark, uncomplicated form, I will, in what follows, present and defend, at least in broad outline, the correct solution to the puzzle AI confronts us with. And AI does present us with a puzzle, because, for reasons we'll investigate in later sections, each of its premises is initially plausible, when H is well chosen. For however improbable or even bizarre it may seem to suppose that I am a BIV, it also seems that I don't know that I'm not one. How *could* I know such a thing? And it also seems that if, for all I know, I am a BIV, then I don't know that I have hands. How could I know that I have hands if, for all I know, I'm bodiless (and therefore handless)? But, at the same time, it initially seems that I do know that I have hands. So two plausible premises yield a conclusion whose negation we also find plausible. So something plausible has to go. But what? And equally importantly, how?

To be sure, the premises are only plausible, not compelling. Thus, we will always have recourse to the Moorean reaction to this

pared to a one-premise argument: I don't know that not-H; so, I don't know that O. The second, "bridge" premise has been added to facilitate my treatment of the argument, nicely dividing those issues that impact on the acceptability of the first premise from those germane to the second.

AI is the first and great argument by skeptical hypothesis. And the second, like unto it, is *The Argument from Possibility* (AP), which, like AI, takes its name from its first premise, and which has this form:

1. It is possible that H_{ind} .
2. If it is possible that H_{ind} , then it is possible that not- O_{ind} .

So, 3. It is possible that not- O_{ind} .

4. If it is possible that not- O_{ind} , then I don't know that O.

So, C. I don't know that O.

(The subscript 'ind' indicates that what occurs in the scope of 'It is possible that' is to be kept in the indicative mood, so that the possibility expressed will be an epistemic one. The "bridge" premises, 2 and 4, can be omitted.) In this paper I address only AI, but let me quickly indicate how AP should be handled. Premise 4, which initially strikes many as AP's weakest link, is actually correct (DeRose 1991, section G). Thus, the AP skeptic must be stopped *before* she reaches step 3. Fortunately, the treatment of AI that I present in this paper can be generalized to handle the initial phase (steps 1–3) of AP as well. This treatment of AP is left here as an exercise for the reader, but is explained in chapter 3, especially section K, of my 1990.

argument: Declare that it is more certain that one knows that one has hands than it is that either of the premises of the argument is true (much less that their conjunction is true), and therefore reject one of those premises, rather than accept the conclusion. But also available is the skeptical reaction, which is to accept the conclusion.

But we should hope for a better treatment of the argument than simply choosing which of the three individually plausible propositions—the two premises and the negation of the conclusion—seems least certain and rejecting it on the grounds that the other two are true. In seeking a solution to this puzzle, we should seek an explanation of how we fell into this skeptical trap in the first place, and not settle for making a simple choice among three distasteful ways out of the trap. We must explain how two premises that together yield a conclusion we find so incredible can themselves seem so plausible to us. Only with such an explanation in place can we proceed with confidence and with understanding to free ourselves from the trap.

Many of those working on AI in recent years seem to have understood this.⁵ And I have good news to report: Substantial progress towards finally solving this skeptical puzzle has been made along two quite different fronts. The bad news is that, as I shall argue, neither approach has solved the puzzle. But the culminating good news is that, as I will also argue, the new solution I present here, which incorporates important aspects of each of the two approaches, *can* finally solve this perennially thorny philosophical problem. While more details and precision will be called for in the resulting solution than I will provide, there will be enough meat on the bones to make it plausible that the fully articulated solution lies in the direction I point to here.

In sections 2–4 of this paper, I explore the contextualist approach to the problem of skepticism, and show why it has thus far fallen short of solving the puzzle. In sections 5–9, I turn to Robert Nozick's attempt to solve our puzzle. Since the shortcomings of Nozick's treatment of knowledge and skepticism have been, at least

⁵This is especially true of Stewart Cohen, to whom I'm indebted for his general setup of the puzzle as a conflict of intuitions, a satisfactory solution of which requires an explanation of why the puzzle arises. See Cohen 1988, 93–94.

to my satisfaction, duly demonstrated by others, it will not be my purpose here to rehearse those shortcomings, but rather to explore and expand upon the substantial insight that remains intact in Nozick's account. In sections 10–17, I present and defend my own contextualist solution, which I argue is the best solution to our puzzle. Since, as I argue in sections 15–17, the skeptic's own solution, according to which we accept AI's conclusion, is among the solutions inferior to the one I present, AI does not successfully support that conclusion.

2. Contextualist Solutions: The Basic Strategy

Suppose a speaker A (for “attributor”) says, “S knows that P,” of a subject S's true belief that P. According to contextualist theories of knowledge attributions, how strong an epistemic position S must be in with respect to P for A's assertion to be true can vary according to features of A's conversational context.⁶

Contextualist theories of knowledge attributions have almost invariably been developed with an eye toward providing some kind of answer to philosophical skepticism. For skeptical arguments like AI threaten to show, not only that we fail to meet very high requirements for knowledge of interest only to misguided philosophers seeking absolute certainty, but that we don't meet even the truth conditions of ordinary, out-on-the-street knowledge attributions. They thus threaten to establish the startling result that we never, or almost never, truthfully ascribe knowledge to ourselves or to other mere mortals.

But, according to contextualists, the skeptic, in presenting her argument, manipulates the semantic standards for knowledge, thereby creating a context in which she can *truthfully* say that we know nothing or very little.⁷ Once the standards have been so

⁶For a bit more on the nature of contextualist theories, see my 1992. The notion of (comparative) strength of epistemic position, central to my characterization of contextualism, will be explicated below in sections 10 and 11.

For exemplary contextualist treatments of the problem of skepticism, in addition to the papers cited below in sections 3 and 4, see especially Unger 1986 and Cohen 1988.

⁷This is at least so according to *skeptic-friendly* versions of contextualist solutions, as will be explained later in this section.

raised, we *correctly* sense that we only could falsely claim to know such things as that we have hands. Why then are we puzzled? Why don't we simply accept the skeptic's conclusion and henceforth refrain from ascribing such knowledge to ourselves or others? Because, the contextualist continues, we also realize this: As soon as we find ourselves in more ordinary conversational contexts, it will not only be true for us to claim to know the very things that the skeptic now denies we know, but it will also be wrong for us to deny that we know these things. But then, isn't the skeptic's present denial equally false? And wouldn't it be equally true for us now, in the skeptic's presence, to claim to know?

What we fail to realize, according to the contextualist solution, is that the skeptic's present denials that we know various things are perfectly compatible with our ordinary claims to know those very propositions. Once we realize this, we can see how both the skeptic's denials of knowledge and our ordinary attributions of knowledge can be correct.

Thus, it is hoped, our ordinary claims to know can be safeguarded from the apparently powerful attack of the skeptic, while, at the same time, the persuasiveness of the skeptical argument is explained. For the fact that the skeptic can invoke very high standards that we don't live up to has no tendency to show that we don't satisfy the more relaxed standards that are in place in more ordinary conversations and debates.

Three important points about contextualist strategies as described above should be made before I move on. First, this type of strategy will leave untouched the timid skeptic who purports by AI merely to be establishing the weak claim that in some (perhaps "high" or "philosophical") sense (perhaps induced by the presentation of AI) we don't know the relevant O, while not even purporting to establish the bold thesis that our ordinary claims to know that same proposition are false. Whether such a timid skeptical stance is of any interest is a topic for another paper. The contextualist strategy is important because AI initially seems to threaten the truth of our ordinary claims—it threatens to boldly show that we've been wrong all along in thinking and saying that we know this and that. For it doesn't seem as if it's just in some "high" or "philosophical" sense that AI's premises are true: They seem true in the ordinary sense of 'know'. In fact, one is initially tempted to say that there's *no* good sense in which I know that I'm

not a BIV or in which I can know I have hands if I don't know that I'm not a BIV. How (and whether) to avoid the bold skeptical result is puzzle enough.

Second, in presenting the contextualist strategy, I have above assumed a skeptic-friendly version of contextualism—one according to which the philosophical skeptic can (fairly easily), and does, succeed in raising the standards for knowledge in such a way as to make her denials of knowledge true. Some contextualists may think that it's not so easy to so raise the standards for knowledge, and that a determined opponent of the skeptic can, by not letting the skeptic get away with raising them, keep the standards low. But the important point is to identify the mechanism by which the skeptic at least threatens to raise the standards for knowledge. Whether the skeptic actually succeeds against a determined opponent in so raising the standards is of little importance. To safeguard ordinary claims to know while at the same time explaining the persuasiveness of the skeptical arguments (which is the goal of his strategy), the contextualist can provisionally *assume* a skeptic-friendly version of contextualism, leaving it as an open question whether and under which conditions the skeptic actually succeeds at raising the standards. The contextualist's ultimate point will then be this: To the extent that the skeptic *does* succeed, she does so only by raising the standards for knowledge, and so the success of her argument has no tendency to show that our ordinary claims to know are in any way defective.

Third, AI can be puzzling even when one is not in the presence of a skeptic who is presenting it. The argument has about the same degree of intuitive appeal when one is just considering it by oneself, without anybody's *saying* anything. But the contextualist explanation, as described above, involves the standards for knowledge being changed by what's being said in a conversation.⁸ For the most part, I will frame the contextualist explanation in terms of such conversational rules, largely because that's what been done by my contextualist predecessors, with whom I want to make contact. But we must realize that the resulting solution will have to be generalized to explain why the argument can be so appealing even when one is considering it in solitude, with nothing being said. The basic idea of the generalization will take either or both of the following

⁸Thanks to Richard Grandy and to Peter Unger for pressing this point.

two forms. First, it can be maintained that there is a rule for the changing of the standards for knowledge that governs the truth conditions of our *thoughts* regarding what is and is not known that mirrors the rule for the truth conditions of what is *said* regarding knowledge. In that case, an analogue of the contextualist solution can be given for thought, according to which the premises and conclusion of AI are truly thought, but my true thought that, say, I don't know that I have hands, had when in the grip of AI, will be compatible with my thought, made in another context, that I do know that very thing. Second, our judgment regarding whether something can or cannot be truly asserted (under appropriate conditions) might be held to affect our judgment regarding whether it's true or false, even when we make this judgment in solitude, with nothing being said at all. That the premises of AI could be truly asserted, then, makes them (at least) seem true even when they're just being thought.

My own solution will employ the basic contextualist strategy explained in this section. But, as should be apparent already, we haven't explained the persuasiveness of AI, and thus haven't solved our puzzle, if we haven't located and explained the conversational rule or mechanism by which the skeptic raises (or threatens to raise) the standards for knowledge. And here contextualists have had little to offer. The two main proposals that have been put forward are discussed in the following two sections.

3. Some Old Contextualist Solutions: Lewis's "Rule of Accommodation"

Though substantial papers have been largely devoted to contextualism and its ability to explain the workings of skeptical arguments like AI, one of the best attempts to explain how (by what rule or conversational mechanism) skeptics raise the standards for knowledge is to be found in David Lewis's "Scorekeeping in a Language Game" (1979), a paper that, while not primarily about knowledge attributions, does treat them in passing.⁹

⁹I am here distinguishing among contextualist solutions according to the mechanism or rule that they allege raises the standards for knowledge. Although there are suggestions of the Relevant Alternatives (RA) approach in "Scorekeeping," Lewis's Rule of Accommodation is quite different from the mechanism most RA theorists posit—thus the separate treatment of

According to Lewis, “rules of accommodation” operate in many spheres of discourse that contain context-sensitive terms.¹⁰ Such rules specify that when a statement is made containing such a term, then—*ceteris paribus* and within certain limits—the “conversational score” tends to change, if need be, so as to make that statement true. For example, ‘flat’, according to Lewis, is a context-sensitive term: how flat a surface must be in order for a sentence describing it as “flat” to be true is a variable matter that is determined by conversational context. And one way to change the conversational score with respect to the standards in place for flatness is to say something that would require for its truth such a change in standards. Suppose, for example, that in a certain conversation the standards for flatness are relaxed enough that my desktop counts as being flat. If I were then to say, “My desktop is not flat,” what I say would be false if it were evaluated according to the standards for flatness in place immediately before this is said. But the Rule of Accommodation specifies that in such a situation—at least under the right circumstances, where the *ceteris paribus* clause is met—the standards for flatness are raised so as to make my statement true.

Lewis suggests that skeptics manipulate a similar rule to change the standards for what is to count as knowledge. According to Lewis’s explanation of the plausibility of skepticism, then, the skeptic’s statements change the conversational score—here, raise the standards for knowledge¹¹—so as to make the skeptic’s statements true. Once the standards for knowledge have been so raised, then

the commonsensical epistemologist must concede defeat. And yet he was in no way wrong when he laid claim to infallible knowledge. What he said was true with respect to the score as it then was.¹² (355)

Lewis. To the extent that Lewis is a relevant alternativist, the RA aspects of his treatment are addressed below in section 4.

¹⁰See especially 346–47.

¹¹For Lewis, as for Relevant Alternatives theorists (see section 4, below), this raising of epistemic standards consists in expanding the range of relevant alternatives to what one believes, that is, the range of alternatives that one must be in a position to eliminate in order to count as knowing.

¹²Why can’t the commonsensical epistemologist simply declare again that he knows, and rely on a Rule of Accommodation to lower the standards back down so as to make *his* claim true? To this Lewis responds that, for some admittedly unknown reason, the standards are more easily raised than lowered (355).

Here Lewis displays the basic contextualist strategy: He protects the truth of what we ordinarily say, or say before the skeptic gets a hold of us, from the skeptic's attack by explaining the success of that attack in terms of the skeptic's changing what counts as knowledge, or, here, "infallible knowledge." Thus, the persuasiveness of the skeptic's attack is explained in such a way as to make it unthreatening to our ordinary claims of knowledge.

And this explanation initially appears to be tailor-made for AI, for AI's first premise is a denial of knowledge—precisely the type of assertion that a rise in the standards for knowledge can help to make true. Such a denial, then, is just the sort of thing that can raise the standards for knowledge via a Rule of Accommodation. Perhaps when the skeptic asserts this first premise, the standards for knowledge are raised, via the Rule of Accommodation, to a level at which we count as knowing neither that we're not BIVs, nor that we have hands.¹³

But a Rule of Accommodation cannot really explain the persuasiveness of AI, or, more generally, of any argument by skeptical hypothesis. To vividly illustrate why this is so, let us imagine and compare two skeptics who are trying to convince you that you don't know that you have hands. The "AI skeptic," true to her name, relies on AI, which, as I noted in section 1, is pretty powerful. The "simple skeptic," on the other hand, simply insists that you don't know that you have hands, offering no reasoning at all for this skeptical assertion.

In seeking a solution to the puzzle generated by AI, we should hope for a solution that, at the very least, explains why the AI skeptic is more convincing than the simple skeptic. If our explanation does not do this much, then we haven't explained how the skeptical argument works on us in any way sufficient to differentiate it from a bald (and dogmatic!) skeptical assertion.

But the Rule of Accommodation, as it stands, appears to be equally accommodating to both of our imagined skeptics. When

¹³To be fair, Lewis, as I've pointed out, treats knowledge only in passing. Although the skeptic he imagines does utilize a skeptical hypothesis (that one is the victim of a deceiving demon (355)), *suggesting* that the treatment Lewis offers should be helpful in solving the puzzle of skeptical hypotheses, he never explicitly attempts a solution to our puzzle. Still, since the solution at least suggested by Lewis is one of the best on offer, it's worth establishing that it can't really solve the puzzle.

the simple skeptic claims that I don't know that I have hands, the supposed Rule of Accommodation should raise the standards for knowledge to such a point as to make her claim true. Of course, the *ceteris paribus* clause may block this result, depending on how it is fleshed out. But there is nothing to this Rule, at least as it has so far been articulated, that would favor the AI skeptic over the simple skeptic. Thus, the explanation based on this Rule does not differentiate between these two skeptics. But if it doesn't do that, it doesn't solve our puzzle.

To avoid possible misunderstanding, let me clearly state that my objection is not to the proposed solution's lack of precision—that we're not given a very clear idea of when the Rule of Accommodation takes effect, that the Rule says merely that the standards *tend* to change in a certain way provided that the (highly unarticulated) *ceteris paribus* clause is met. My own solution will be likewise imprecise. No, the problem isn't that the Rule isn't completely filled in, but rather that, for the reasons given above, since the explanatory work needed to solve the puzzle isn't done by the aspects of the Rule that have been provided, it will have to be done by just those aspects of the Rule that haven't been provided. And, as we've little idea what these aspects are, we've little idea of what it is that may solve the puzzle.¹⁴

Perhaps, when it's more fully articulated, the operative Rule of Accommodation can be seen to contain a feature that favors the AI skeptic over the simple skeptic. In that case, the solution to our puzzle, which has so far eluded us, may (at least in part) be found in a fuller articulation of that Rule.

But I doubt that the solution even lies in that direction. One

¹⁴None of this is to deny that there is some Rule of Accommodation according to which the standards for knowledge tend to be raised to "accommodate" denials of knowledge. Nor is it even to deny that such Rules of Accommodation help the AI skeptic. In fact, I find it plausible to suppose that many denials of knowledge, including those of AI skeptics, often do exert an upward pressure on the standards for knowledge via some such rule. Likewise, certain settings (in addition to courts of law, certain philosophy classes are good examples), it seems to me, tend to militate in favor of high epistemic standards. AI skeptics may take advantage of these factors, the influence of which may explain some of the persuasiveness of their skeptical performances. But to solve our puzzle, we want primarily to explain what the nature of the skeptical argument itself adds to the effectiveness of the skeptic's performance that goes beyond what is contributed by the skeptic's setting and the fact that she asserts her conclusion.

(secondary) reason for my doubt is that positive claims to know that skeptical hypotheses don't obtain seem to raise the standards for knowledge as well as do denials of such knowledge.

To illustrate this I'll use Fred Dretske's familiar example of mules cleverly painted to look like zebras (Dretske 1970, 1015–16). If I saw what looked to be zebras in the zebra cage at a zoo, I would ordinarily claim to know that the animals in the cage are zebras. (Suppose, for instance, that my son asked me, "Do you know what those animals are?" I would respond positively.) A skeptic might challenge this supposed knowledge with an instance of AI where O is *Those animals are zebras* and H is *Those animals are mules cleverly painted to look like zebras*. The resulting premises are individually plausible, since I couldn't tell a cleverly painted mule from a zebra. A contextualist treatment of this instance of AI will claim that in asserting the first premise, the skeptic raises the standards for knowledge to a level at which I count as knowing neither that the animals are not cleverly painted mules nor that they're zebras.

And it indeed does seem that once this skeptical hypothesis is brought into play, I cannot happily claim to know what I so happily claimed to know before. To be in a good enough position to claim to know that the animals are zebras according to the standards brought into play by the skeptic, one must be in a good enough position that one can rule out¹⁵ the hypothesis that they are cleverly painted mules. Since I'm not in that kind of epistemic position, I don't count as knowing, although perhaps someone more familiar with mules and zebras would still count as knowing, even at these higher standards—someone, for instance, who was in a position to say, "No, they can't be mules: no mule's head is shaped like that."

But these same higher standards seem to be induced when the skeptical hypothesis is brought into play by a positive claim to know that it doesn't obtain. Suppose, to vary Dretske's example, that I am confronted, not by a skeptic, but by a boastful zoologist. He brags, "Due to my vast knowledge of zebra and mule anatomy, I know that those animals are not mules cleverly painted to look like zebras; so I know that they're really zebras." This zoologist, as much as the skeptic, seems to invoke higher standards for knowl-

¹⁵For some comments on this notion of "ruling out" see sections 4 and 5, below.

edge at which he, but not I, will count as knowing that the animals are zebras. He certainly seems to be claiming more than the mundane knowledge that even I possess—and claim to possess—in an ordinary zoo setting, where there's no such zoologist telling me what's what.

But a Rule of Accommodation cannot account for *this* rise in standards, for the zoologist doesn't deny any supposed knowledge. To the contrary, what he does is make positive claims to know, and a rise in standards for knowledge can never help to make true a positive claim to know. So, as I said, a Rule of Accommodation can't do anything to explain this notable rise in epistemic standards.¹⁶

My primary reason for doubting that our solution is to be found in a fuller articulation of the Rule of Accommodation is this: To explain the persuasiveness of AI (and, in particular, of its first premise) in such a way as to differentiate the AI skeptic from the simple skeptic, we must identify the feature of skeptical hypotheses that makes it particularly hard to claim or to think that one knows that *they* are false. Far from being found in a Rule of Accommodation, then, a solution to our puzzle, if it's to be found at all, is to be found in an explanation of what it is about skeptical hypotheses that makes these propositions, as opposed to ever so many other propositions, such effective skeptical weapons. So, to solve the puzzle, we must locate or articulate this peculiarly potent feature of just these propositions (the skeptical hypotheses). And, once we see what this feature is and how it works, the Rule of Accommodation is destined to play only a rather subsidiary role (see note 14) in explaining the effectiveness of the skeptic's attack.

My secondary reason for doubting that the Rule of Accommodation might solve our puzzle was worth bringing up both because

¹⁶It's been proposed to me, on behalf of the Rule of Accommodation and the solution to AI that can be based on it, that the boastful zoologist, while he does not say that I don't know, does strongly suggest or imply that I don't, and the Rule of Accommodation operates here on his suggestion: the standards go up so as to make the suggestion true. I am skeptical of this attempt to salvage the solution for two reasons. First, I suspect that the rule becomes far too powerful if it's allowed to work on what we suggest as well as on what we say. Second, the standards for knowledge seem likewise raised even if the boastful zoologist thinks I am also an expert, and thinks he is informing me that he too knows what's what. Here he's not even suggesting that I don't know.

it seems to me to have some force, and because it vividly illustrates this important fact: The upward pressure on the standards for knowledge that bringing skeptical hypotheses into play exerts is exerted whether the hypotheses are raised in denials of knowledge or in positive claims to know.

4. Some Old Contextualist Solutions: The “Relevant Alternatives” Approach and the Rule of Relevance¹⁷

Perhaps the most popular solution to our puzzle has been put forward by advocates of the “Relevant Alternatives” theory of knowledge (RA). Again suppose a speaker A says, “S knows that P.” According to RA, such an assertion is made within and must be evaluated against a certain framework of *relevant alternatives* to P. To know that P is to have a true belief that P and to be able to rule out these relevant alternatives. But not every contrary of or alternative to P is a *relevant* alternative.¹⁸ In an ordinary case of claiming to know that some animals in the zoo are zebras, to again use Dretske’s example, the alternative that they’re cleverly painted mules is not relevant. Thus, I can truthfully claim to know they’re zebras despite my inability to rule out this fanciful alternative.

But in various *extraordinary* cases, the painted mules hypothesis is a relevant alternative. It might be made relevant by some extraordinary feature of S (the putative subject of knowledge) or her surroundings.¹⁹ But most RA theorists are contextualists, and allow

¹⁷Fred Dretske (see his 1970, 1971, 1981a, 1981b), although he does advocate a Relevant Alternatives theory of knowledge, proposes a treatment of AI quite different from that described below. I’m not certain whether Dretske’s is even a contextualist version of RA. (As I note in part 2 of my 1992, one can be an RA theorist without being a contextualist.) One thing is clear about Dretske’s treatment of AI: He denies premise (2). Given this, his treatment runs into the same difficulties as does Nozick’s; see especially section 9 below.

¹⁸See, for example, Dretske 1970, 1022; Goldman 1976, 772; and Stine 1976, 249.

¹⁹Thus, if S is at a zoo that fairly consistently uses painted mules in an attempt to fool the zoo-going public, then the painted mule hypothesis is relevant. So, even though S is lucky enough to be at this zoo on one of the rare days when actual zebras are being used, S cannot truthfully be said to know that they’re zebras unless she is able to rule out the painted mule hypothesis, which she can’t do unless she knows more than I do about zebras and mules.

that features of the conversational context in which A (the ascriber of knowledge) finds himself, in addition to features of S and her surroundings, can influence which alternatives are relevant.²⁰ Alvin Goldman, for instance, suggests that “if the speaker is in a class in which Descartes’s evil demon has just been discussed,” then certain alternatives may be relevant that ordinarily are not (1976, 776).

It is this contextualist aspect of (most versions of) RA that facilitates the most commonly proposed solution to our puzzle, the Relevant Alternatives Solution (henceforth, ‘RAS’). With some slight variations in detail in different presentations of it, the basic idea of RAS is this: The AI skeptic’s *mentioning* of the BIV hypothesis in presenting the first premise of AI *makes* that hypothesis relevant. Once the skeptical hypothesis has been made relevant, we correctly sense that we cannot truthfully claim to know anything contrary to it unless we can rule it out. Since we are unable to rule it out, and since it is an alternative to both *I am not a BIV* and to *I have hands*, we correctly sense that we could only falsely claim to know these things. So the skeptic truthfully asserts that we don’t know that the hypothesis doesn’t obtain, and then truthfully concludes that we don’t know that we have hands.²¹

Why then are we puzzled? Because we at the same time realize that the BIV hypothesis is not ordinarily relevant. We realize that in most of the conversational circumstances in which we find ourselves, our inability to rule out the skeptic’s far-fetched hypothesis is no bar to our truthfully claiming to know such things as that we

²⁰As I explain in part 2 of my 1992, an RA theorist can be an invariantist if he allows only factors about the putative subject of knowledge and her surroundings, and not conversational factors pertaining to the speaker (the ascriber of knowledge), to affect which alternatives are relevant. Matters get tricky with first-person knowledge claims, where S and A are identical. Here, in addition to allowing features that affect how good an epistemic position our subject actually is in, and that thereby attach to her qua putative subject of knowledge, contextualist RA theorists will also allow features of her conversational context, which affect how good a position she must be in to count as knowing, and which thereby attach to her qua attributor of knowledge, to influence what the range of relevant alternatives is.

²¹Again, here I’m only giving the skeptic-friendly version of this contextualist solution. An RA theorist might be less friendly to the skeptic by holding, for example, that mentioning an alternative makes that alternative relevant only if one’s conversational partner lets one get away with making it relevant.

have hands. Thus, even as we find the skeptic's denials of knowledge persuasive, we realize that when we again find ourselves in more ordinary contexts, it will not only be correct for us to claim to know such things, it would be wrong to deny that we know them merely because we can't rule out the BIV hypothesis. What we fail to realize, according to RAS, is that our ordinary claims to know such things as that we have hands are compatible with the skeptic's present denial that we know those very things.

RAS, then, is an instance of the general contextualist strategy—one according to which the raising of the standards consists in enlarging the range of alternatives that are relevant and that one must therefore be in a position to rule out in order to count as knowing. The conversational rule or mechanism that RAS posits for enlarging that range (raising the standards for knowledge), then, is that *mentioning* a proposition Q—*ceteris paribus* and within certain limits, no doubt—tends to make Q a contextually relevant alternative to any P that is contrary to Q. Call this the *Rule of Relevance*.²²

Note that this Rule of Relevance, as opposed to the Rule of Accommodation, can handle cases like that of the boastful zoologist, in which a positive claim to know that a skeptical hypothesis doesn't obtain seems to have the same effect on the meaning of sentences containing 'know' as would a denial of such knowledge. This is to be expected on the present Rule of Relevance, on which both the denial and the claim to know will, by including a mention of the skeptical hypothesis, expand the range of relevant alternatives so that it will include that hard-to-rule-out hypothesis.

²²Of course, it shouldn't be held that just any mention of a proposition makes that proposition a relevant alternative. In order to be made relevant, the proposition must, no doubt, be inserted into a conversation in the right way. But the advocate of RAS can plausibly claim to have explained the persuasiveness of AI even if he hasn't given an exact specification of the conditions under which a mentioning of a proposition makes that proposition a relevant alternative. Plausibly holding that in presenting AI the skeptic *does* insert her skeptical hypothesis into the conversation in the right way, the advocate of RAS can leave it as a future project to specify more exactly just which ways are the right ways. Although this by itself will be neither necessary nor sufficient for the mentioning of a proposition to be of the right kind to enlarge the range of relevant alternatives so as to include it, it nonetheless may be relevant that in the skeptic's presentation of AI's first premise, the mentioning of the hypothesis occurs within the scope of an epistemic operator—"S does not know that. . ."

But to explain the persuasiveness of AI (particularly of its first premise), and to thereby solve our puzzle, a treatment of AI must tell us what it is about skeptical hypotheses that makes it difficult to claim to know that they don't obtain. The key feature of skeptical hypotheses that RAS seizes on is clearly this: we can't rule them out.

And isn't there something to this explanation? For it seems that we indeed can't rule out (effective) skeptical hypotheses, and it further seems that it is precisely this fact that makes them such effective skeptical weapons.

But though it is plausible to suppose we can't rule out skeptical hypotheses, and also plausible to say that we don't know that they don't obtain, it is futile to try to explain the plausibility of the latter by that of the former.

Indeed, there are plenty of other phrases that can be used plausibly to describe our apparently limited epistemic position with regard to effective skeptical hypotheses. All of the following descriptions about my position vis-à-vis the BIV hypothesis have some initial plausibility: I cannot rule it out, I don't know that it doesn't obtain (and don't know whether it obtains), I can't tell that it doesn't obtain (and can't tell whether it obtains), I can't discern that it doesn't obtain (and can't discern whether it obtains), and I can't distinguish its obtaining from its not obtaining, and so on, and so forth. But citing one of these to explain the plausibility of another doesn't occasion even the slightest advance in our understanding.

What accounts for the plausibility of saying that I don't know that I'm not a BIV? The fact that I can't discern that I'm not one? This is no explanation. It seems just as good (in fact, to me, better) to reverse things and claim that the fact that I don't know that I'm not a BIV accounts for the plausibility of saying that I can't discern that I'm not one.

Likewise for ruling out. It is indeed plausible to suppose that we can't rule out skeptical hypotheses. And it's plausible that we don't know that they don't obtain. But it doesn't seem to advance our understanding much to explain the plausibility of either by that of the other.

(An exercise for the reader: Randomly pick two of the above negative assessments of our epistemic position vis-à-vis effective skeptical hypotheses. Then consider whether the plausibility of the

first can be explained by reference to the second. Then reverse things and consider whether the plausibility of the second can be explained by reference to the first. Try the same procedure on another pair of descriptions. (If you're running low on such negative assessments, you'll find it's easy, following my lead, to come up with many more on your own.) Then evaluate the success of explaining the plausibility of AI's first premise by reference to the fact that we can't rule out effective skeptical hypotheses.)

To explain why we feel some pull towards describing our epistemic position with regard to skeptical hypotheses in any of the above less than flattering ways—as well as very many other ways that I didn't bother to mention—we need an explanation that reaches outside this circle of all-too-closely related terms of epistemic appraisal.²³ Indeed, as will emerge in the following sections (especially section 8), the best explanation for the plausibility of AI's first premise also seems to provide a good account of why it seems that we can't rule out skeptical hypotheses, as well as an explanation of the plausibility of the various other pessimistic evaluations. Once this explanation is in place, it becomes even clearer that none of the things it's used to explain can be properly used to explain each other.

5. The Subjunctive Conditionals Account (SCA) of the Plausibility of AI's First Premise

The main stumbling block of the contextualist solutions we've discussed has been a failure to explain what it is about skeptical hypotheses that makes it so plausible to suppose that we don't know that they're false. This point of weakness in the contextualist solutions is the particular point of strength of Nozick's treatment of AI in his *Philosophical Explanations* (1981). In this and the following three sections I'll present and defend the *Subjunctive Conditionals*

²³Goldman (1976) cashes out “discriminating” what one believes from a relevant alternative to it in terms of what one would believe if the alternative obtained. This, combined with the Rule of Relevance, could yield an approach to skepticism close to the one I'll here defend. Goldman himself does not propose a solution to the skeptical problem; he strives to remain neutral on the issue. But I'll be working in the general direction I think Goldman points to.

Account (SCA) of the plausibility of AI's first premise, which I've abstracted from Nozick's account of knowledge and skepticism.

According to SCA, the problem with my belief that I'm not a BIV—and I do have such a belief, as do most of us—is that I would have this belief (that I'm not a BIV) even if it were false (even if I were one). It is this that makes it hard to claim to *know* that I'm not a BIV. For, according to SCA, we have a very strong general, though not exceptionless, inclination to think that we don't know that P when we think that our belief that P is a belief we would hold even if P were false. Let's say that S's belief that P is *insensitive* if S would believe that P if P were false. SCA's generalization can then be restated as follows: We tend to judge that S doesn't know that P when we think S's belief that P is insensitive.

As is well worth noting, this general inclination explains the operation of nonphilosophical skeptical hypotheses that are far less radical than the BIV hypothesis or even the painted mule hypothesis. Just so, it serves to explain why, even though I feel inclined to say that I know the Bulls won their game last night because I read the result in a single newspaper, I still feel strongly pulled toward admitting the (mildly) skeptical claim that I don't know that the paper isn't mistaken about which team won: I realize that my belief that the paper isn't mistaken is a belief I would hold even if it were false (even if the paper were mistaken).

Indeed, after encountering a couple of instances of AI with different skeptical hypotheses plugged into the 'H' slot (for example, the BIV, the painted mules, and the mistaken paper hypotheses), one develops a sense of what makes for an effective skeptical hypothesis and, thus, an ability to construct convincing instances of AI oneself. To make AI's second premise convincing, it is usually sufficient (though not necessary) that H be incompatible with O. But what about the first premise? To make *it* convincing, we instinctively look for a hypothesis that elicits in the listener both the belief that the hypothesis doesn't obtain and an acknowledgement that this belief is one she would hold even if the hypothesis *did* obtain.

Upon hearing the hypothesis, typically one can't help but projecting oneself into it. How would things seem to me if that situation obtained? Well, pretty much (or sometimes exactly) as they actually seem to me. And, so, what would I believe if such a "strange" situation obtained? Pretty much (or exactly) what I ac-

tually believe. For example, and in particular, if I *were* a BIV, I would believe every bit as firmly as I actually do that I *wasn't* one. But if this belief is one I would hold even if it were false, how can I be in a position to tell that, or discern that, or *know* that, it's true?

As I've just hinted, a similar explanation, in terms of subjunctive conditionals, can explain the plausibility of the other ways we feel inclined to describe our seemingly limited epistemic position vis-à-vis effective skeptical hypotheses. Consider especially the description involving 'ruling out'. In a normal zoo setting, most of us would take ourselves to know that the animals in the zebra cage are zebras. From this, it seems, we should be able to infer that they're not cleverly painted mules, since zebras aren't mules. So why are we reluctant to count our seeing the zebras and performing this inference as a case of ruling out the painted mule hypothesis? Because, the explanation goes, even after performing the inference, it still seems we would believe the observed animals weren't painted mules if they were precisely that. Why does it seem we can't tell that they're not painted mules? Because we would believe they weren't even if they were. Ditto for why we seemingly can't discern that they're not and why it seems we can't distinguish their being cleverly painted mules from their not being such, etc.

Also worth noting is the usefulness of SCA in explaining our reluctance to ascribe knowledge in certain lottery situations. Even where the odds of your being a loser are astronomically high (there are 20 million tickets, only one of which is a winner, and you have but one ticket), it can seem that you don't know that you're a loser of a fair lottery if the winner hasn't yet been announced. SCA accounts for this seeming: Your belief that you're a loser is one you would hold even if you were the winner.

SCA is a powerful explanation. But there are problems. As I suggested above, there are exceptions to the general inclination to which SCA appeals: There are cases in which it seems to us that some S does know that P even though we judge that S would believe that P even if P were false. Some of these exceptions will be quickly discussed in sections 6 and 7 below. The first and main point to make regarding such exceptions, of course, is that this very general inclination needn't be exceptionless to perform the explanatory role SCA assigns it. In section 8 we will see strong grounds for endorsing SCA as being at least on the right track despite the exceptions to the generalization to which it appeals. But these exceptions

are still worth examining, for they will indicate certain important directions in which SCA can be improved, even though we won't be in a position to make SCA ideally precise here.

6. SCA, Grandmothers, and Methods

First, then, consider a case discussed by Nozick:

A grandmother sees her grandson is well when he comes to visit; but if he were sick or dead, others would tell her he was well to spare her upset. Yet this does not mean she doesn't know he is well (or at least ambulatory) when she sees him. (1981, 179)

Here, it seems, the grandmother knows her grandson is well, though it can seem that she doesn't satisfy the third condition of a preliminary form of Nozick's analysis of S knows that P, which is:

- (3) If p weren't true, S wouldn't believe that p.

Nozick's response is to relativize this third condition to the method by which S has come to believe that p, yielding:

- (3) If p weren't true and S were to use M to arrive at a belief whether (or not) p, then S wouldn't believe, via M, that p (179),

where 'M' is the method by which S has come to believe that p.²⁴

Unlike Nozick, I'm not presenting an analysis of propositional knowledge. But his grandmother case also seems to be an exception to the general inclination SCA appeals to: Here we're not at all inclined to think the grandmother doesn't know her grandson is well, even though it can seem that if he weren't well, she would still believe he was. The generalization SCA utilizes says that we tend to judge that S doesn't know where S does not satisfy Nozick's

²⁴Precisely, what Nozick does is this: He analyzes the technical locution 'S knows, via method M, that p', and then in turn analyzes the relation of S's knowing that p in terms of this technical locution. The revised third condition I've displayed is part of Nozick's attempt to analyze the technical locution.

third condition for knowledge. One possibility here is to follow Nozick *very* closely by modifying that generalization so that it refers to Nozick's modified, rather than his original, third condition, and thus, like Nozick, explicitly relativizing our account to the method by which S believes that P.

Often, though, context takes care of this for us. Even to one aware of the likelihood that the grandmother's family would have kept her in the dark about her grandson's condition were he not well, it *can* seem that even Nozick's initial formulation of the third condition for knowledge is met by the grandmother. On one way of evaluating that simple conditional, it seems that if the grandson were not well, the grandmother would *not* believe he was well. After all, she's looking right at him! The standard possible-worlds semantics for counterfactual conditionals can illuminate what's going on here. When one searches for the possible worlds most similar to the actual world in which the grandson is not well, the respects in which the possible worlds are to resemble the actual world is a highly context-sensitive matter. Especially where the context focuses one's attention on the grandmother and her cognitive and recognitional abilities, one *can* place heavy weight upon similarity with respect to the method she is using to arrive at her belief, and then it can seem that in the closest world in which the grandson is not well, she's looking right at him and seeing that he's not well, and so does *not* believe he is well. On this way of evaluating the conditional, the grandmother *does* satisfy even the initial formulation of Nozick's third condition, and she's no counter-example to the generalization utilized by SCA. But, in evaluating that simple conditional, one can also stress other similarities, particularly ones involving the propensities and plans of the various family members (or whatever facts ground the judgment that if her grandson weren't well, the grandmother would be effectively lied to), to reach the verdict that if he were not well, she *would* believe that he was well.

We can sharpen SCA by specifying that we tend to judge that S doesn't know when she fails to satisfy Nozick's initial formulation of (3), where (3) is evaluated in such a way that heavy emphasis is put upon similarity with respect to the method of belief formation utilized by S, or, following Nozick, we can insert a specification of the method into the antecedent of (3).²⁵ But in neither case is this

²⁵These are not identical modifications. On the first option, similarity with respect to method is weighted heavily, but can be outweighed by other

to make a very precise modification; rather, it merely indicates the direction in which a more precise account might lie, for any such use of the notion of *methods* of belief formation in our account invites a host of questions (many of which Nozick wrestles with) involving how such methods are to be specified and individuated.

7. SCA and Some Skeptical Hypotheses That Don't Work

Certain instances of AI aren't very persuasive. The first premise of the argument can be quite unconvincing despite the fact that SCA predicts that we'd find it plausible. Suppose, for instance, that in an attempt to show by AI that I don't know I have hands, a skeptic utilizes, instead of the BIV hypothesis, the following simple H: I falsely believe that I have hands. The resulting instance of AI seems to pack little or no more punch than a simple skeptic's unsupported claim that I don't know I have hands. It's at the first premise that this ill-fated instance of AI fizzles. But my belief that I don't falsely believe that I have hands is insensitive: If this belief were false (if I did falsely believe that I have hands) I would still believe it was true (I'd still believe that I don't falsely believe that I have hands). Likewise insensitive is my belief that the following hypothesis is false: I'm an intelligent dog who's always incorrectly thinking that I have hands. If this belief of mine were false (if I were such a deluded intelligent dog) I'd still believe it was true (I'd still believe that I wasn't such a creature). So SCA, as it has so far been formulated, predicts that it will seem to us that the above beliefs don't amount to knowledge and that we'll find plausible the first premise of AI that results when the above hypotheses are used. But in fact these instances of AI's first premise are far from convincing. As opposed to the BIV hypothesis, it seems that one *does* know that the deluded dog hypothesis and the simple false belief hypothesis are false.

Again, the main point to make here is that SCA's generalization

factors. Thus, even so evaluated, the most similar world(s) in which the antecedent of the original (3) are true may be worlds that diverge from the actual world with respect to the method by which S came to believe that P. By contrast, on the second option, since the method by which S believes that P becomes part of the antecedent of the conditional we're evaluating (the modified (3)), the closest possible world(s) in which that antecedent is true cannot be worlds that diverge from the actual world with respect to method.

needn't be exceptionless to be explanatory. While a more precisely Chisholmed refinement of SCA might not have the negations of these ineffective H's as instances of those propositions it says we tend to judge we don't know, I'll here just make a preliminary observation as to what might be going wrong. Part of the problem with these "hypotheses" is that they don't give us much of an idea of *how* I come to have the false belief they assign to me. Hypotheses are supposed to explain; skeptical hypotheses should explain how we might come to believe something despite its being false. The first of these hypotheses simply stipulates that I'm wrong about my having hands, without indicating how I came to be so sadly mistaken. The second adds to the first that I'm a dog, which adds little to our understanding of how my mistake about having hands came about. By contrast, when we encounter effective skeptical hypotheses, we have some understanding of how (if H is true) we have come to falsely believe that O. If either of our ineffective hypotheses is filled in so as to make it clear to us how I came to falsely believe I have hands, it becomes effective.

SCA's generalization was this: We tend to judge that S doesn't know that P when we think that S's belief that P is insensitive (when we think that S would believe P even if P were false). The limitation of SCA's generalization that's suggested by these cases is this: We *don't* so judge ourselves ignorant of P where not-P implies something we take ourselves to know to be false, without providing an explanation of how we came to falsely believe this thing we think we know. Thus, *I falsely believe that I have hands* implies that I don't have hands. Since I do take myself to know that I have hands (*this* belief isn't insensitive), and since the above italicized proposition doesn't explain how I went wrong with respect to my having hands, I'll judge that I do know that proposition to be false. But this again is just a preliminary statement, and there's room for a lot more refinement here. What we need now is some assurance that we're headed in the right direction.

8. SCA Confirmed

Such assurance is to be found by considering what it would take to make it seem to us that we *do* know skeptical hypotheses to be false.

But let's first reconsider the lottery case. As noted above in section 5, we are puzzlingly reluctant to claim knowledge in certain lottery situations. The explanation provided by SCA for this phenomenon is intuitively appealing: It does seem that the fact that we would believe that we were losers even if we were winners is largely what's behind our judgment that we don't know we're losers. SCA receives further powerful support when we consider the grounds that *do* seem to us sufficient for knowledge of one's being a loser. In the lottery situation, even a very minute chance of being wrong seems to deprive one of knowledge. But if we're going to worry about even such minute chances of error, then why does it seem that you do know you're a loser after the winning number has been announced on the radio and you've compared the numbers on your ticket with the sadly different numbers announced? After all, radio announcements *can* be in error; what you're hearing *may* not be a real radio announcement but the voice of a friend who's rigged up a practical joke; you *might* be suffering from some weird momentary visual illusion and misreading the numbers on your ticket; and so forth. All very remote possibilities, to be sure. But, since we're already countenancing even the most minute chances of error, why don't these possibilities rob us of knowledge even after the announcement has been made and heard?

SCA's explanation of why we don't think we know *before* the announcement is made is that we at that time judge that if we weren't losers, we'd still believe that we were. Note that once you've heard the announcement of the winning numbers and compared them with the numbers on your ticket, it *no longer* seems that if you had been the winner, you'd believe you were a loser. Rather, we judge that in that case you'd now believe you were the winner or would at least be suspending judgment as you frantically double-checked the match. It's very impressive that the very occurrence that would suffice to make it seem to us that you do know you're a loser (the radio announcement) also reverses our judgment regarding the truth of the conditional appealed to in SCA to explain why it seems to us that you don't know before the announcement is made. The occurrence which gets us to judge that we know here also removes what SCA posits as the block to our judging that we know. This is an indication that SCA has correctly identified the block.

SCA similarly provides a very intuitively appealing explanation for why it seems to us that we don't know that skeptical hypotheses

are false, as was also noted in section 5. It again receives powerful further confirmation as we look to cases in which one seemingly does know that a skeptical hypothesis doesn't obtain (cases in which skeptical hypotheses that are ordinarily effective fail to be effective). The boastful zoologist I introduced toward the end of section 3, it seems, knows that the animals in the zebra cage are not cleverly painted mules, while I, it seems, do not. But the very anatomical knowledge that seemingly enables him to know they're not painted mules also has the consequence that if the animals *were* cleverly painted mules, the zoologist, unlike me, would *not* believe that they weren't. And although I don't seem to know they're not painted mules simply by looking at them, I could, it seems, get to know this if I undertook some special investigation—perhaps, as has been suggested in the literature (Stine 1976, 252), one involving paint remover. Which special investigations would do the trick (and under which circumstances would they)? A survey of various scenarios yields an impressive correlation: The investigations that would seemingly allow me to know that the animals aren't painted mules would also affect our judgment as to the truth value of the subjunctive conditional so critical to SCA. Once I have completed the investigation, it seems that I, like the zoologist, would *not* believe that the animals weren't painted mules if in fact they were. Likewise, by checking appropriately independent sources, I could get myself into a position in which I seemingly *would* know that the newspaper isn't mistaken about whether the Bulls won last night. But the checks that would seemingly allow this knowledge would also make it seem that if the paper were mistaken, I would *not* believe it wasn't. Again and again, SCA posits a certain block to our judging that we know, and the changes that would clear the way for our judging that we know also remove this block. This makes it difficult not to believe that SCA is at least roughly correct.

In the case of the BIV hypothesis, it's hard to test SCA in this way, for it's difficult to imagine a situation in which it seems a subject does know that she's not a BIV. But this only confirms SCA: While it's difficult to imagine a situation in which one seems to know that one's not a BIV, it's likewise difficult to imagine circumstances in which the block SCA posits is removed. It's difficult, that is, to imagine a situation in which someone believes they're not a BIV but in which the conditional *If S were a BIV, then S would believe*

she wasn't a BIV isn't true. For, as the BIV hypothesis is formulated, one's brain is electrochemically stimulated so that one has precisely those sensory experiences one actually has had. But wouldn't one then have formed precisely those beliefs that one actually has formed, including the belief that one's not a BIV?

Unlike that involved in the Relevant Alternatives Solution, the present explanation for the plausibility of AI's first premise can't be happily reversed: Trying to account for the plausibility of the subjunctive conditional *If H were true, I would believe it was false*, by reference to the (presumed) fact that I don't know that H is false certainly seems to get things backwards. Much better to follow the proposed Nozickean route in explaining the plausibility of denying knowledge by reference to the conditional.

Further, as was the case with not knowing, the investigations that would reverse our other pessimistic judgments regarding your standing vis-à-vis a skeptical hypothesis would also put you in a position to say that you wouldn't believe the hypothesis is false if it were true. Thus, for instance, to make it seem that you can tell that those animals aren't painted mules, you must put yourself in such a position that you wouldn't believe they weren't if they were. And, as was the case with not knowing, none of these explanations by subjunctive conditionals seems happily reversible.

It seems that this explanation, SCA, for the plausibility of AI's first premise must be (at least roughly) correct and, therefore, that it points to part of the solution to our puzzle.

Indeed, some readers will wonder why I have claimed only that our general tendency not to count insensitive beliefs as instances of knowledge explains that premise's plausibility and have stopped short of accepting sensitivity as a necessary condition for knowledge²⁶ and therefore simply endorsing that first premise as true. But while we've just seen strong grounds for simply accepting AI's first premise, there are also strong grounds for accepting AI's second premise and for accepting the denial of its conclusion. We have to stop short somewhere; we can't simply accept all three members of this triad as true. To solve this puzzle, I'll claim that

²⁶Or, given the exceptions to the general tendency that we've discussed in sections 6 and 7, why I haven't accepted that some properly Chisholmed refinement of the sensitivity requirement (which has as instances of it convincing instances of AI's first premise) is necessary for knowledge.

AI's first premise, while not *simply* true, is true according to unusually high standards for knowledge. But, I'll argue, my solution explains why that premise seems true and, more generally, why sensitivity seems necessary for knowledge. If my solution provides the best explanation for how all three members of our puzzling triad seem true, that will be good reason for stopping short where my solution tells us to, rather than where one of its inferior rivals—bold skepticism, for example—tells us to.

9. Nozick's Own Solution and the Abominable Conjunction

Nozick's own treatment of AI, from which SCA was abstracted, fails. This treatment is based on Nozick's account of knowledge as true, *sensitive* belief, where, very roughly, one's true belief that *p* is sensitive to the truth value of *p* if one would not have believed that *p* if *p* had been false.²⁷ Thus, Nozick's treatment of AI involves accepting the skeptic's first premise. But, at the same time, and much more unfortunately, it also involves denying the second. You *don't* know that you're not a BIV, Nozick claims, because any belief you might have to this effect is insensitive: You would have held this belief even if it were false (even if you were a BIV). By contrast, Nozick claims, your belief that you have hands *is* a sensitive belief: If *it* were false—if you didn't have hands—you would not hold it. So you do know you have hands even though you don't know that you're not a BIV. The skeptic's mistake—the second premise—is supposing that you can know you have hands only if you also know that you're not a BIV.

Or so Nozick claims. This is not the place for a general evaluation of Nozick's analysis of propositional knowledge, so let us confine ourselves to the results of this analysis as applied to the beliefs in question in AI. Here Nozick's account does very well in issuing the intuitively correct verdict for the relevant particular judgments regarding what is known and what is not. Most of us would judge

²⁷Though this statement of Nozick's account of knowledge is rough, that will not affect my treatment, which would apply equally well to Nozick's full account. I've skipped entirely Nozick's fourth condition for knowledge, but I believe this fourth condition to be redundant, anyway: It automatically holds whenever true belief is present. Also, as I've already noted, Nozick takes account of the method of belief formation in his final version of the third condition. The same thing happens with the fourth.

that we do know such things as that we have hands, and this is Nozick's verdict. And, when a skeptical hypothesis is well chosen, it does seem quite plausible to most of us that we don't know that it doesn't obtain. But there are three relevant issues to our puzzle: Is the first premise of AI true? Is the second premise true? Is the conclusion true? And it's easy to endorse the intuitively correct answer to two out of the three questions if you're willing to take the implausible stand on the remaining one.

Nozick takes his implausible stand on the issue of the second premise, denying it in the face of its evident intuitive appeal.²⁸ Accepting his treatment involves embracing the abominable conjunction that while you don't know you're not a bodiless (and handless!) BIV, still, you know you have hands. Thus, while his account does quite well on the relevant particular intuitions regarding what is and isn't known, it yields an intuitively bizarre result on the comparative judgment the second premise embodies.²⁹

²⁸At 205–6 Nozick admits this appeal, and later he writes, “Thus, if our notion of knowledge was as strong as we naturally tend to think (namely, closed under known logical implication) then the skeptic would be right. (But why do we naturally think this? Further exploration and explanation is needed of the intuitive roots of the natural assumption that knowledge is closed under known logical implication)” (242).

Nozick is quite hard on anti-skeptics who choose rather to deny the first premise; he writes: “The skeptic asserts we do not know his possibilities don't obtain, and he is right. Attempts to avoid skepticism by claiming we do know these things are bound to fail. The skeptic's possibilities make us uneasy because, as we deeply realize, we do not know they don't obtain; it is not surprising that attempts to show we do know these things leave us suspicious, strike us even as bad faith” (201). But similar remarks could be made about Nozick. As Nozick himself admits, the second premise has its own intuitive appeal. So why not say that what we “deeply realize” is that if you don't know that you're not a BIV, then you don't know you have hands, and that the skeptic is right about *this*? Nozick's denial of the second premise leaves me about as “suspicious” as does a denial of the first, and though Nozick's denial doesn't strike me as an instance of bad faith, denials of the first premise seem no better candidates for that charge.

²⁹What are Nozick's grounds for rejecting the second premise? Nozick notes that the premise is an instance of a very general principle to the effect that knowledge is closed under known implication (see note 33, below). After admitting that the closure principle *seems* true (205–6), Nozick claims that it's wrong, and his reasons for this claim are made entirely from within his analysis of knowledge: Given his analysis, knowledge won't be closed (see especially 206–8). So Nozick is relying on his analysis to show us that the second premise is false despite its intuitive appeal. And indeed, Nozick has developed and defended his analysis of knowledge (in

As promised, I won't here rehearse the powerful objections to Nozick's analysis of propositional knowledge that have been put forward,³⁰ but, assuming that this analysis isn't independently convincing before we turn to the problem of skeptical hypotheses,³¹ we're left with little reason to follow Nozick in choosing to take an implausible stand precisely where he has rather than someplace else.

This leaves us in a bind. For, as we saw in sections 5 and 8 above, SCA is quite powerful. That explanation is that we realize that any belief we might have to the effect that an (effective) skeptical hypothesis doesn't obtain is insensitive, and we're inclined to think that insensitive beliefs don't constitute knowledge. How can we appropriate that explanation without following Nozick in having to implausibly deny the second premise of AI and embrace the abominable conjunction?

10. Strength of Epistemic Position and AI's Second Premise

Here's how: by incorporating SCA into a contextualist solution to our puzzle that avoids such a fumbling of AI's second premise. Indeed, I propose a *very* strong endorsement of that second premise.

Recall that according to contextualist theories of knowledge attributions, how strong a subject's epistemic position must be to make true a speaker's attribution of knowledge to that subject is a flexible matter that can vary according to features of the speaker's conversational context. Central to contextualism, then, is the notion of (*relative*) *strength of epistemic position*. In presenting and defending contextualism, I've found that most listeners feel that they understand pretty well what's meant when I claim, for instance,

part 1 of chapter 3) before he applies it to the issue of skepticism (in part 2).

³⁰Unfortunately, what is perhaps the most powerful attack on Nozick's theory of knowledge, made by Saul Kripke in lectures, circa 1985, has not, to the best of my knowledge, found its way into print. For those interested in critical literature on Nozick, a good place to start is with Forbes 1984 and several of the essays in Luper-Foy 1987. For still further reading, Luper-Foy 1987 contains an excellent bibliography.

³¹As remarked in note 29, Nozick depends heavily on the independent plausibility of this analysis to provide the momentum for his treatment of AI.

that sometimes the standards for knowledge are higher than usual, or that in some conversational situations one's epistemic position must be stronger than in others to count as knowing. But it would be good to clarify this important notion of strength of epistemic position as best we can by, for instance, supplying an intuitive test for when one epistemic position is stronger than another. The best such device is that of *comparative conditionals*. One can have a variety of grounds for assenting to conditionals like *If Mugsy is tall, then Wilt is tall*, and *If Wilt is not tall, then Mugsy is not tall*. But one very good basis for assenting to these conditionals is the comparative knowledge that Wilt is at least as tall as Mugsy. Likewise, where S is a putative subject of knowledge, P is a true proposition that S believes, and A and B are situations in which S is found, we can have similarly comparative grounds for assenting to conditionals of the form *If S knows that P in A, then S knows that P in B*. In such a case, the comparative grounds for our assent is our realization that S is in *at least as strong* an epistemic position with respect to P in situation B as he is in with respect to that same proposition in situation A, and this comparative conditional serves as a good intuitive test for that comparative fact: It brings that fact to light.

So, for instance, to borrow some examples from Alvin Goldman (1976), let Henry be our subject, and let *What Henry is seeing is a barn* be the thing Henry putatively knows. Both in situation F (for "fakes") and in situation N ("no fakes"), Henry is driving through the countryside and, having no reason to think there's anything unusual going on, very firmly believes, and takes himself to know, that the object he's seeing is a barn. And indeed, in both cases, it is a barn. But in F, unbeknownst to him, Henry is in an area that is filled with very convincing fake barns—papier-mâché barn facades. In fact, we may suppose that Henry has just been fooled more than twenty times by such fakes, although he's now looking at the only actual barn for miles around, and so this time truly believes that what he's seeing is a barn. N is exactly like F, except that there are no fakes in the area—the things Henry has taken to be barns have all actually been barns. With regard to these examples, the conditional *If Henry knows in F, then he knows in N* seems to get the comparison right, indicating that Henry's in at least as strong an epistemic position in situation N as he is in situation F. The evident failure of *If Henry knows in N, then he knows in F* to get the comparison right shows that Henry's not in as strong a position

to know in F as in N. Together, these results indicate that Henry's in a stronger epistemic position in N than in F.

As is important to our discussion of AI's second premise, comparative conditionals can similarly be used to test the relative strength of epistemic position of a single subject with respect to *different propositions* that subject believes in the same situation: Thus, the intuitive correctness of *If S knows that P, then S knows that Q* and *If S doesn't know that Q, then S doesn't know that P* can indicate that S is in at least as strong an epistemic position with respect to Q as she's in with respect to P.³²

Sometimes no clear verdict results when we attempt to evaluate a conditional in this comparative way, for the good reason that it's unclear how the two epistemic positions we're evaluating compare with one another. Thus, if we compare a situation in which Henry has a good look at the barn but in which there are a couple of fake barns several miles away that Henry hasn't encountered with a situation in which there are no fakes at all in Henry's vicinity but in which he doesn't have quite as good a look at the barn, the relevant conditionals can be difficult to evaluate. But, in many instances, some of the relevant conditionals *are* clearly true on comparative grounds.

Such is the case with instances of AI's second premise, where the skeptical hypothesis is well chosen. They seem true and *are* true, I suggest, for just this comparative reason: As we realize, we *are* in at least as good a position to know that the hypothesis is false as we're in to know the targeted piece of presumed ordinary knowledge.³³ Let's look briefly at some instances. Recall the following epistemologically perplexing pairs of propositions:

³²And, of course, such conditionals can be used to make all manner of other comparisons: comparative strength of the epistemic positions of two *different subjects* with respect to the same proposition or with respect to different propositions, the strength of the epistemic position of a subject with respect to one proposition in one situation as compared with that same subject's epistemic position with respect to a different proposition in a different situation, etc.

³³As is well known, instances of AI's second premise are often instances of the principle that knowledge is closed under known logical implication: $Kp \ \& \ K(p \text{ entails } q) \rightarrow Kq$. (In the next paragraph I explain why this is not always the case, at least when the closure principle isn't strengthened as there described.) As is also well known, there are exceptions to the principle so formulated, and it might take a lot of tinkering to get it exactly right. But, as Nozick, the arch denier of closure, puts it, "We would be ill-

not-H

I'm not a BIV.

Those animals aren't just cleverly
painted mules.The paper isn't mistaken about
whether the Bulls won last night.O

I have hands.

Those animals are
zebras.The Bulls won last
night.

Given natural background assumptions, we can sense that the following comparative fact holds for each of the above pairs: I am in no better a position to know that O than I am in to know that not-H. This comparative fact is revealed in each case by the highly plausible conditional that is AI's second premise: If I don't know that not-H, then I don't know that O. Closely tied to that comparative fact in each case is the related and intuitively compelling realization that it would be no wiser to bet one's immortal soul on O's being true than to bet it on not-H's being true.

I propose then to accept the relevant conditional with respect to each of the above pairs, and to accept other convincing in-

advised, however, to quibble over the details of P [the principle that knowledge is closed under known logical implication]. Although these details are difficult to get straight, it will continue to appear that something like P is correct" (1981, 205). Nozick goes on to claim that this appearance is deceiving. I believe that something like P is correct, but that doesn't compete with my present account of AI's second premise: When a conditional is an instance of the properly formulated closure principle, the relevant comparative fact involving strength of epistemic position holds. See Brueckner 1985 for arguments that the denial of knowledge closure principles "is not a fruitful anti-skeptical project" (112).

While restrictions will have to be put on the closure principle that will weaken it in certain respects, there may be other respects in which it can be strengthened. Some instances of AI's second premise are convincing even though H is compatible with O. For instance, the BIV hypothesis seems to undermine my putative knowledge of *I'm in Houston* as well as of *I have hands*, but, of course, that I'm a bodiless BIV is compatible with my being in Houston. Perhaps if S is to know that P, then S must know that not-Q for any Q (but here restrictions must be added) such that if Q were true, S would not know that P. Thus, the range of Qs that must be known not to obtain may be broadened so as to include not only propositions that are incompatible with P, but also others such that if they were the case, then S wouldn't know that P. Those Qs that *are* incompatible with P itself will then be seen as special cases of those that are at odds with S's knowing that P. Barry Stroud discusses a stronger closure principle such as this in his 1984 (25–30).

stances of AI's second premise. Indeed, these conditionals are true *regardless of how high or low the standards for knowledge are set*. Just as the comparative fact that Wilt is at least as tall as Mugsy has the result that the conditional *If Wilt is not tall, then Mugsy is not tall* will be true regardless of how high or low the standards for tallness are set, so the comparative fact that I'm in at least as strong an epistemic position with respect to not-H as I'm in with respect to O will result in *If I don't know that not-H, then I don't know that O* being true regardless of how high or low the standards for knowledge are set. Thus, we will never have to follow Nozick in accepting the abominable conjunction: that conjunction is false at any epistemic standard.

With that ringing endorsement of AI's second premise anchored firmly in place, we can return to the first premise, hoping to incorporate SCA into a contextualist account of that premise's plausibility.

11. Strength and Sensitivity

As has become very apparent, two notions that are central to my attempt to solve our puzzle are, on the one hand, the Nozickean notion of the sensitivity of beliefs and, on the other, the notion of strength of epistemic position. While both notions stand in need of a good deal of sharpening and explanation (only some of which they'll receive here), we've already obtained interesting results applying them to the epistemologically perplexing pairs of propositions displayed above. In each case, one's belief in O is sensitive, while one's belief in not-H is insensitive. Yet, at the same time, one is in at least as strong an epistemic position with respect to not-H as one is in with respect to O.

For each of the second and third pairs of propositions, one could gather further evidence, strengthen one's epistemic position with respect to both not-H and O, and *make* even one's belief that not-H sensitive. But even before this further evidence is gathered, one's belief that O is *already* sensitive, despite the fact that one is in no stronger an epistemic position with respect to this O than one is in with respect to not-H. (With respect to the first pair of propositions, it is difficult to imagine a situation in which one is in such a strong position with respect to one's not being a BIV that this belief is sensitive.)

This leads us to an important insight regarding skeptical hypotheses: One's epistemic position with respect to propositions to the effect that skeptical hypotheses don't hold must be stronger than it is with respect to other, more ordinary propositions (e.g., our above Os) if belief in such propositions is to be sensitive.

An explanation of our two central notions in terms of possible worlds will provide a partial and quite rough-and-ready, but still somewhat enlightening, picture of how this situation can arise. An important component of being in a strong epistemic position with respect to P is to have one's belief as to whether P is true match the fact of the matter as to whether P is true, not only in the actual world, but also at the worlds sufficiently close to the actual world. That is, one's belief should not only be true, but should be non-accidentally true, where this requires one's belief as to whether P is true to match the fact of the matter at nearby worlds. The further away one can get from the actual world, while still having it be the case that one's belief matches the fact at worlds that far away and closer, the stronger a position one is in with respect to P. (Recalling the results of section 6, we should remember either to restrict our attention solely to those worlds in which the subject uses the same method of belief-formation she uses in the actual world, or to weigh similarity with respect to the subject's method very heavily in determining the closeness of possible worlds to the actual world.) If the truth-tracking of one's belief as to whether P extends far enough from actuality to reach the closest not-P worlds, then one doesn't believe that P in those closest not-P worlds, and one's belief that P is sensitive. But how far from actuality must truth-tracking reach—how strong an epistemic position must one be in—to make one's belief that P sensitive? That, of course, depends on how distant from actuality the closest not-P worlds are.

Consider my belief that I have hands. I believe this at the actual world, and it's true. What's more, in the other nearby worlds in which I have hands, I believe that I do. There are also, at least in my own case, some alarmingly close worlds in which I don't have hands. These include worlds in which I lost my hands years ago while working on my uncle's garbage truck. In the closest of these not-P worlds, I'm now fully aware of the fact that I'm handless, and my belief as to whether I have hands matches the fact of the matter. My belief as to whether I have hands doesn't match the fact in various worlds in which I'm a BIV, of course, but these are *very*

distant. While there are closer worlds in which the match fails, it seems that in a fairly wide range of worlds surrounding the actual world, my belief as to whether I have hands does a good job of matching the fact of the matter. Thus, I'm in a pretty strong epistemic position with respect to that matter.

Now let P be *I'm not a BIV*. Where not-P (here, *I am a BIV*) is quite remote, one can be in a quite strong epistemic position with respect to P merely by believing that P in all the nearby worlds. As I do believe this P in such nearby worlds, I'm in a pretty strong epistemic position with respect to this P. This can occur, and in my case, does occur, even though one's belief as to whether P doesn't match the fact of the matter in the closest not-P worlds: Since even the closest of the not-P worlds are quite distant, one's belief as to whether P needn't match the fact of the matter that far from the actual world for one to be in a quite strong position with respect to P.

But for one's belief that P to be sensitive, one must *not* believe that P in the closest not-P worlds. Since skeptical hypotheses tend to fasten on somewhat remote (and sometimes very remote) possibilities, then, one can be in a relatively (and sometimes a very) strong position with respect to beliefs to the effect that they don't obtain (since one's belief as to whether they obtain matches the fact of the matter over a wide range of worlds closest to the actual world), while these beliefs remain insensitive (since one would still believe that the hypotheses didn't obtain in the closest worlds in which they do obtain). By contrast, where P is such that there are both P and not-P worlds very close to the actual world, one's belief that P must be sensitive (one must not believe that P in the closest not-P worlds) in order for one to be in even a minimally strong epistemic position with respect to P, and, conversely, one needn't be in a very strong position for one's belief to be sensitive.

12. The Rule of Sensitivity and the Beginnings of a New Contextualist Solution

The important insight regarding skeptical hypotheses—that one's epistemic position with respect to propositions to the effect that skeptical hypotheses don't hold must be stronger than it is with respect to other propositions before beliefs in such propositions can

be sensitive—suggests a new contextualist account of how, in presenting AI, the skeptic raises the standards for knowledge. Let's call the conversational rule this new account posits as the mechanism by which the skeptic raises the standards for knowledge the "Rule of Sensitivity." Although a more general formulation of this rule is desirable, I will here state it in such a way that it applies only to attributions (and denials) of knowledge, since such applications are what's needed to address the present puzzle.³⁴ So limited, our rule is simply this: When it is asserted that some subject S knows (or does not know) some proposition P, the standards for knowledge (the standards for how good an epistemic position one must be in to count as knowing) tend to be raised, if need be, to such a level as to require S's belief in that particular P to be sensitive for it to count as knowledge. Where the P involved is to the effect that a skeptical hypothesis does not obtain, then this rule dictates that the standards will be raised to a quite high level, for, as we've seen, one must be in a stronger epistemic position with respect to a proposition stating that a skeptical hypothesis is false—relative to other, more ordinary, propositions—before a belief in such a proposition can be sensitive.

A story in terms of possible worlds again provides a rough-and-ready, but still perhaps enlightening, picture of how the Rule of Sensitivity operates. Context, I've said, determines how strong an epistemic position one must be in to count as knowing. Picture this requirement as a contextually determined sphere of possible worlds, centered on the actual world, within which a subject's belief as to whether P is true must match the fact of the matter in order for the subject to count as knowing. (Given the results of section 6, we must again remember either to restrict our attention solely to those worlds in which the subject uses the same method of belief formation she uses in the actual world, or to weigh similarity with

³⁴Introducing a skeptical hypothesis into a conversation in any number of ways other than in attributions and denials of knowledge can seem to raise the standards for knowledge. For instance, instead of arguing, "You don't know that the paper isn't mistaken about the result of last night's game; therefore, you don't know that the Bulls won," a skeptic may urge, "Consider this proposition: The newspaper is mistaken about who won the game. Now, keeping that proposition clearly in mind, answer me this: Do you *really* know that the Bulls won?" Of course, as with the Rule of Relevance (see note 22), not just *any* mention of a skeptical hypothesis seems to trigger the mechanism for raising the standards of knowledge I'm about to articulate.

respect to the subject's method very heavily in determining the closeness of possible worlds to the actual world.) Call this sphere the sphere of epistemically relevant worlds. As the standards for knowledge go up, the sphere of epistemically relevant worlds becomes larger—the truth-tracking of one's belief must extend further from actuality for one to count as knowing. Given this picture, the Rule of Sensitivity can be formulated as follows: When it's asserted that S knows (or doesn't know) that P, then, if necessary, enlarge the sphere of epistemically relevant worlds so that it at least includes the closest worlds in which P is false.

A powerful solution to our puzzle results when we follow the basic contextualist strategy (see section 2) and utilize this Rule of Sensitivity to explain how the standards for knowledge are raised by the skeptic's presentation of AI. While many noteworthy features and virtues of this solution are best explained by comparing it with the other proposed solutions to our puzzle, as I'll do in following sections, the basic idea of the present solution is this. In utilizing AI to attack our putative knowledge of O, the skeptic instinctively chooses her skeptical hypothesis, H, so that it will have these two features: (1) We will be in at least as strong a position to know that not-H as we're in to know that O, but (2) Any belief we might have to the effect that not-H will be an insensitive belief (a belief we would hold even if not-H were false—that is, even if H were true). Given feature (2), the skeptic's assertion that we don't know that not-H, by the Rule of Sensitivity, drives the standards for knowledge up to such a point as to make that assertion true. By the Rule of Sensitivity, recall, the standards for knowledge are raised to such a level as to require our belief that not-H to be sensitive before it can count as knowledge. Since our belief that not-H isn't sensitive (feature (2)), the standards are driven up to such a level that we don't count as knowing that not-H. And since we're in no stronger an epistemic position with respect to O than we're in with respect to not-H (feature (1)), then, at the high standards put in place by the skeptic's assertion of AI's first premise, we also fail to know that O. At these high standards, the skeptic truthfully asserts her second premise (which, recall, is also true at lower standards), and then truthfully asserts AI's conclusion that we don't know that O.³⁵ This

³⁵Again, I'm here assuming a skeptic-friendly version of contextualism. See the second important point made at the end of section 2.

accounts for the persuasiveness of AI. But since, on this account, the skeptic gets to truthfully state her conclusion only by raising the standards for knowledge, AI doesn't threaten the truth of our ordinary claims to know the very Os our knowledge of which the skeptic attacks. For the fact that the skeptic can install very high standards that we don't live up to has no tendency to show that we don't satisfy the more relaxed standards that are in place in more ordinary conversations and debates.

13. The Rule of Sensitivity and SCA: A Comparison of Our New Solution with the Other Contextualist Solutions and with Nozick's Solution

Recall that the problem with the other contextualist solutions we've seen is that they fail to adequately explain why AI's first premise has the intuitive pull it has (when the skeptical hypothesis employed is well chosen). Our new contextualist solution gains an important advantage over its contextualist rivals by incorporating SCA. We explain the plausibility of AI's first premise by reference to the following two facts. First, any belief we might have to the effect that a skeptical hypothesis doesn't obtain (where that hypothesis is well chosen) is insensitive: as we realize, we would hold this belief even if it were false (even if the hypothesis did obtain). And, second, we have a very general inclination to think that we don't know that P when we realize that our belief that P is insensitive—when we realize that we would believe that P even if P were false.

We follow Nozick in employing SCA. But we diverge from Nozick's treatment in our account of why the second fact above holds. On Nozick's account, we have the general inclination asserted there because our concept of knowledge just is, roughly, that of true, sensitive belief. This would account for our inclination to deny the status of knowledge to insensitive beliefs alright, but it would also have us happily asserting abominable conjunctions, which, in fact, we're loathe to do.

Our new solution avoids this unhappiness by not building a sensitivity requirement into the very concept of knowledge. The notion of sensitivity, rather, finds its happier home in our contextualist account of how the standards for knowledge are raised, and

the second fact above is accounted for as follows. Where S's belief that P is not sensitive, S is not in a good enough epistemic position to count as knowing that P by the standards that, according to the Rule of Sensitivity, would be put in place by the very claim that S knows (or doesn't know) that P. Thus, an assertion that S doesn't know that P, where S's belief that P is insensitive, will raise the standards for knowledge to a level high enough to make that denial of knowledge true. A positive claim that S *does* know such a P, on the other hand, is doomed to failure: The making of the claim will raise the standards for knowledge to a level high enough to make that claim false. So, whenever S's belief that P is insensitive, we can *truthfully* assert that S *doesn't* know that P, and can only *falsely* say that S *does* know that P. No wonder, then, that the second fact holds!

Thus, we successfully incorporate SCA, explaining the plausibility of AI's first premise, without following Nozick in licensing abominable conjunctions.

14. Our New Contextualist Solution Clarified and Compared with the Straightforward Solution

The puzzle of skeptical hypotheses, recall, concerns the premises of AI together with the negation of its conclusion:

1. I don't know that not-H.
2. If I don't that not-H, then I don't know that O.

not-C. I do know that O.

A solution to the puzzle must, of course, issue a verdict as to the truth of each of these three, but it must also explain why we find all of them plausible.

Let's be clear about what our present contextualist solution has to say about each of these. Our verdict regarding (2) is that it's true regardless of what epistemic standard it's evaluated at, so its plausibility is easily accounted for. But this, combined with a similarly enthusiastic endorsement of (1), would land us in bold skepticism. We avoid that fate by endorsing (1) as true, not at all standards, but only at the unusually inflated standards conducive to skepticism. Thus, on our solution, we do know, for instance, that we're not BIVs, according to ordinary low standards for knowledge.

But, though (1) is false when evaluated according to those ordinary low standards, we're able to explain its plausibility, as we've seen, by means of the fact that the high standards at which (1) is true are precisely the standards that an assertion or denial of it put into play. Since attempts to assert (1) are bound to result in truth, and attempts to deny it are destined to produce falsehood,³⁶ it's no surprise that we find it so plausible.

But what of (not-C)? On the present solution, claims to know ordinary propositions are true according to ordinary low standards but false according to the highly inflated standards that, by the Rule of Sensitivity, are put in place by the assertion of (1). (Not-C) seems plausible because it's true when evaluated at the standards most normally applied to it. But, it will be asked, why do we find these claims to know plausible even when we're in a context in which the skeptic has raised the standards to such a level that these claims are false? A little caution is in order here. It's controversial just how intuitively correct (not-C) does seem to us in such a context. Most of us feel some ambivalence. Such ambivalence is to be expected whenever we're dealing with a puzzle consisting of mutually inconsistent propositions, all of which are individually plausible. For when the propositions are considered together, one will have this good reason for doubting each of them: that the others seem true. And it's difficult to distinguish the doubt of (not-C) that arises from this very general source (that its falsehood follows from other things one finds plausible) from that which arises from the fact that the standards are high. At any rate, the very strong pull that (not-C) continues to exert on (at least most of) us even when the standards are high is explained in the manner out-

³⁶But for cases in which it seems one *can* truthfully say "S knows that not-H," despite the fact that S's belief that not-H is insensitive, see chapter 3, section J ("Low-Strength Claims to Know that Skeptical Hypotheses Do Not Obtain") of my 1990. In such cases, given certain features of the conversational situation, the Rule of Sensitivity does not operate. These constitute exceptions to the rule that one cannot truthfully call an insensitive belief knowledge. As I explain there, I welcome these exceptions, and would actually be a bit worried if there weren't such exceptions. For it's a feature of my treatment of AI that we do know skeptical hypotheses to be false according to low epistemic standards. I would find it a bit embarrassing if we could never *claim* to have such knowledge by means of simple knowledge attributions, and I'm reassured by the result that in special conversational circumstances, it seems we *can* truthfully claim to know that not-H, despite the fact that our belief that not-H is insensitive.

lined in section 2: Even while we're in a context governed by high standards at which we don't count as knowing that O, we at the same time realize that as soon as we find ourselves in more ordinary conversational contexts, it will not only be true for us to claim to know these very Os that the skeptic now denies we know, but it will also be wrong for us to deny that we know these things. It's easy, then, to think that the skeptic's present denial must be equally false and that it would be equally true for us now, in the skeptic's presence, to claim to know that O.

The verdicts the present solution issues regarding the truth values of the members of the triad are complicated by the fact that ours is a contextualist solution. Only (2) receives the same verdict regardless of what the epistemic standards are; the truth values of (1) and (not-C) vary with context. It's just this variance that our solution so essentially relies on in explaining how we fall into our puzzling conflict of intuitions. Noncontextualist (henceforth, "straightforward") solutions, on the other hand, must choose one of the members of this triad to deny, claiming this loser to be false according to the invariant epistemic standards that govern all attributions and denials of knowledge: The "Moorean" solution in this way denies (1),³⁷ the "Nozickean" (2), and the "Bold Skeptical" solution thus denies (not-C), accepting that we speak falsely whenever, even in ordinary, nonphilosophical discussions, we claim to know the O in question.

From the perspective of our present contextualist solution, each of these straightforward solutions results in part, of course, from a failure to see the truth of contextualism. But which straightforward solution an invariantist confusedly adopts will depend on the standards that dominate her evaluation of our beliefs in O and in not-H. If her evaluation is dominated by the relatively low standards that govern our ordinary, out-on-the-street talk of knowledge, she will end up a Moorean. If she evaluates the beliefs in question according to the high standards that are put into place by the skeptic's presentation of AI, bold skepticism is the result. The Nozickean solution ensues from evaluating each belief according to

³⁷This is called the "Moorean" solution because Moore responded in this way to the dream argument. It's far from certain that Moore would have so responded to other instances of AI that utilize different skeptical hypotheses.

the standards that would most often be used in evaluating that belief. For reasons we've seen, a claim to know (or an admission that one doesn't know) that a skeptical hypothesis is false will, by the Rule of Sensitivity, tend to invite a very high reading, at which the admission is true and the claim is false. But a claim to know that *O* doesn't so demand a high reading. From the present perspective, the Nozickean is reacting to the fact that one can usually truthfully claim that one does know that *O* and can usually truthfully claim not to know that not-*H*. What the Nozickean misses is how difficult it is to make these two claims together: once you have admitted that you don't know that not-*H*, it seems the reverse of intuitively correct to claim to know that *O*, at least until the conversational air is cleared.

To succeed, a straightforward solution must explain what leads our intuitions astray with respect to the unlucky member of the triad which that solution denies. Otherwise, we'll have little reason for denying just that member of the triad. Nozick himself provides no such explanation with respect to (2), parenthetically leaving this vital task to "further exploration,"³⁸ and other Nozickeans, if any there be, have not, to the best of my knowledge, progressed any farther along this front. Mooreans, to the best of my knowledge, have fared no better in explaining why we're so reluctant to claim the status of knowledge for our insensitive beliefs. It's the defenders of bold skepticism who've made the most progress here. In the remaining sections, I'll explain why our contextualist solution is superior to that of the bold skeptic.

15. Bold Skepticism and the Warranted Assertability Maneuver

Almost all of the time, it seems to almost all of us that we do know the *O*s that the skeptic claims we don't know. According to the bold skeptic, whenever we say or think that we know these things, we say or think something false. The bold skeptic thus implicates us, speakers of English, in systematic and widespread falsehood in our use, in speech and in thought, of our very common word 'know'. Equally paradoxically, the bold skeptic holds that we're speaking the truth whenever we say that someone doesn't know

³⁸See the first paragraph of note 28, above.

these Os, even though it seems to most of us that we'd then be saying something quite false. What leads us astray? Peter Unger and Barry Stroud have suggested on behalf of bold skepticism that although we don't know these O's, it's often useful for us to claim that we do know them, and we are therefore often warranted or justified in making such claims. What then leads us astray is this: We mistake this useful/justified/warranted assertability of knowledge ascriptions for truth.³⁹ On the other side of the coin, presumably, we're mistaking the useless/unwarranted/unjustified assertability of denials of knowledge for falsehood.

Two serious problems emerge for the bold skeptic at this point. The first is that such "warranted assertability maneuvers" could be attempted by advocates of the other solutions as well. Warranted assertability indeed can be mistaken for truth, and unwarranted assertability for falsehood, but this by itself does not favor the bold skeptic's solution over the other straightforward approaches. Each of the straightforward approaches denies a member of the triad constituting our puzzle, and each it seems could claim that the reason this loser they've chosen seems true, though it's in fact false, is that we're often warranted in asserting it, and we mistake this warranted assertability for truth. Thus, the Moorean, for instance, could claim that although we do indeed know that H is false, we're not warranted in claiming that we know this (though this claim would be true), but are rather warranted in saying that we don't know (though this latter is false). Simply attributing apparent truth to warranted assertability is a game almost any party to this dispute can fairly easily play.⁴⁰ That this line of thought would eventually work out any better for the bold skeptic than for his opponents would take some showing.⁴¹

It's at (1) that the skeptic has his best hope of gaining an advantage over my solution, for that premise indeed does seem true,

³⁹This is the basic line Unger takes in his defense of bold skepticism in his 1975; see especially pages 50–54. Stroud, though not himself advocating bold skepticism, does seek to defend the bold skeptic along these lines in chapter 2 of his 1984; see especially pages 55–82.

⁴⁰By contrast, our new contextualist solution attributes the apparent truth of (1) to (1)'s *truth* (and not just its warranted assertability) at the very standards its assertion invokes.

⁴¹For my own part, for reasons I can't go into here, I think the resulting Moorean position would be slightly more defensible; thus, if I had to reject contextualism and adopt a straightforward solution, I'd be a Moorean.

and, unlike the skeptic, I've stopped short of fully endorsing it, making do with an explanation of its plausibility. But the skeptic's other problem lurks here. Usually, while solving a philosophical puzzle consisting of a set of individually plausible but mutually inconsistent claims, one only has to explain (away) the plausibility of those members of the set one denies, and one is relieved of the burden of explaining the plausibility of those members that one endorses, their truth and our ability to recognize that truth being explanation enough of their apparent truth. But truth does not suffice to explain apparent truth where one makes us out to be absolutely horrible judges of truths of the kind in question. Thus, the skeptic's second big problem is that, because he holds that we're subject to constant and radical error as to the scope of our knowledge, consistently thinking we know things when we don't, the skeptic, although he thinks (1) is true, owes us an explanation for its plausibility. Given that our habit of mistaking our ignorance for knowledge is so pervasive, why doesn't it seem to us *here* that we know what, in fact, we don't—that these skeptical hypotheses are false? Why does our lack of knowledge, which we're so pervasively blind to, shine through so clearly to us just where the issue is whether we know a skeptical hypothesis to be false?

The skeptic's initial answer will certainly be that we're *not* warranted in claiming to know that skeptical hypotheses don't obtain, and thus can't mistake warranted assertability for truth here. But then, to see why skeptical hypotheses are effective, we must be told why we're not warranted in claiming to know that skeptical hypotheses are false, given that, according to the skeptic, we are warranted in claiming to know all manner of other things that in fact we don't know. And here skeptics have little to offer. But if the results of sections 5 and 8 above are correct, the answer must involve the lack of sensitivity enjoyed by our beliefs that skeptical hypotheses don't obtain. The skeptic's use of SCA will take this form: Although we know nothing (or very little), it's when our beliefs are insensitive that we're not even warranted in asserting that we know and we therefore recognize our lack of knowledge. But the skeptic must now also address AI's second premise, making sure his endorsement of SCA is made in such a way as to account for our intuitions here. Indeed, whether or not he buys into SCA, the skeptic faces this question: If, as he claims, we're usually under the delusion that we know that O, but we customarily recognize

that we don't know that not-H, why aren't we happy to conjoin this error with that insight and embrace the abominable conjunction?

This may look like a difficult question, but the skeptic has a ready answer. His problem is that the warranted assertability maneuver by itself didn't really solve our puzzle, but rather re-introduced it in a new form. And the only way I've seen to incorporate SCA into a treatment of AI that also handles the other pieces of our puzzle is to employ the idea that contextually sensitive epistemic standards govern our use of 'know', and to posit the Rule of Sensitivity as the mechanism by which the AI skeptic drives those standards up, as I've advocated here. But wise invariantists typically accept that contextually varying standards govern our use of ascriptions and denials of knowledge. The sensible invariantist will admit that, of course, what passes for knowledge in some contexts won't so pass in others. Being an invariantist, he'll deny that the truth conditions of knowledge attributions vary in the way the contextualist claims they do. But the clever invariantist will maintain that the varying epistemic standards that the contextualist supposes govern the truth conditions of these sentences in fact govern their conditions of warranted assertability.⁴²

This allows the bold skeptic to mimic any contextualist solution, and in particular the solution I'm advocating here, by means of a simple twist. With respect to my solution, the bold skeptic can maintain that the Rule of Sensitivity is a rule for the raising of the epistemic standards governing our use of sentences ascribing knowledge to subjects, alright, but insist that it governs the warranted assertability conditions of these sentences, rather than their truth conditions, which, he'll maintain, remain constant at a level beyond the reach of mere mortals to satisfy. The warranted assert-

⁴²Stroud thus claims that on the skeptic's conception of our practices, we operate under certain "practical constraints" (1984, 75) in our everyday uses of 'know', and asserts that our standards for saying we know vary from case to case (65–66). Thus, on the skeptic's conception, the standards for ascribing knowledge that we employ in everyday use depend upon our "aims and interests at the moment" (65). According to contextualism, these varying standards reflect a corresponding variation in the truth conditions for attributions of knowledge. But on Stroud's skeptic's conception, when we ascribe knowledge in everyday situations, we are typically saying something literally false, although "the exigencies of action" justify these false attributions. The best exploration of this type of idea is provided by Unger in his 1984.

ability maneuver can then be employed: We mistake warranted assertability for truth (and unwarranted assertability for falsehood). Thus, since we're never warranted in claiming to know that skeptical hypotheses don't obtain (due to the operation of the twisted Rule of Sensitivity), we're led to judge (correctly) that such claims to knowledge would be false. And since AI's second premise is always warranted, we judge (again correctly) that this premise is true. But since a claim to know some O is usually warranted, due to the low standards for warranted assertability that would ordinarily be applied to such a claim, we judge (incorrectly) that we know this O. Thus, my solution, like other contextualist solutions, can be easily adapted to suit the purposes of the bold skeptic. The result is a theory parallel to my own contextualist solution, which differs in its semantics of 'know': According to this parallel invariantist theory, the context-sensitive varying epistemic standards we've discovered govern the warranted assertability conditions of attributions and denials of knowledge, rather than their truth conditions, which are held to be invariant.⁴³ How shall we rationally decide between a contextualist solution, and in particular the one I'm here defending, and the bold skeptic's analogue of it?⁴⁴

16. Bold Skepticism and Systematic Falsehood

Like its contextualist relatives, our new solution is designed largely with the goal in mind of crediting most of our attributions of knowledge with truth. And no wonder. We in general take it as a strike against a theory of a common term of a natural language that it involves the speakers of that language in systematic and widespread falsehood in their use of that term. Let's borrow an example and suppose, for instance, that a crazed philosopher

⁴³Going back to the bold skeptic's first problem, note that all this maneuvering can be mimicked by the Moorean, who can also hold that a Rule of Sensitivity governs the warranted assertability conditions of knowledge ascriptions. Like the bold skeptic, the Moorean can hold that the truth conditions of such attributions of knowledge remain invariant, but in the Moorean's hands, these constant epistemic standards will be meetably low.

⁴⁴Readers of Unger's 1984 will see the strong influence of that excellent book on my procedure here, though I come to very different conclusions than he does in that work. (But see his more recent 1986.)

claimed that there are no physicians, because, in addition to holding a medical degree, a necessary condition for being a physician is that one be able to cure any conceivable illness.⁴⁵ On what grounds should we reject this bizarre conjecture in favor of a more traditional and less demanding account of what it is to be a physician? Our language certainly could have been such that S's having the ability to cure any conceivable illness was a truth condition of 'S is a physician' (although the word 'physician' would not have been very useful in that case). In virtue of what is our language in fact such that the strange theory is not true of it? I'm of course not in a position to give a complete answer to this question, but it's eminently reasonable to suppose that such facts as these, regarding our use, in thought and in speech, of the term 'physician' are involved: that we take to be physicians many licensed practitioners of medicine who don't satisfy the demanding requirement alleged; that we seriously describe these people as being physicians; that we *don't deny* that these people are physicians; etc. It's no doubt largely in virtue of such facts as these that the traditional view, rather than the conjecture of our crazed philosopher, is true of our language. (The correctness of the traditional view largely *consists* in such facts.) And these facts also provide us with our best reasons or evidence for accepting the traditional, rather than the strange, hypothesis regarding the semantics of 'physician'. In this case, that the peculiar theory implicates us in systematic and widespread falsehood in our speech and thought involving 'physicians' is a (constitutive and evidential) strike against the theory that proves quite decisive.

If our crazed philosopher tried to account for the above facts regarding our use of the term 'physician' via the quick and easy conjecture that the less demanding requirements that are more traditionally assigned to 'physician', while they don't accurately specify the truth conditions of sentences involving that term, do articulate these sentences' warranted assertability conditions, we should not, on the basis of this maneuver, suspend our judgment against his contention. That his theory involves us in systematic falsehood continues to constitute a strike against it, and in the absence of quite weighty counterbalancing considerations that fa-

⁴⁵See Stroud (1984, 40), who in turn borrowed the example from elsewhere.

vor the strange theory over the traditional one, this strike remains decisive.

Of course, the problem with this hopeless nonstarter of a theory is that there don't seem to be any such counterbalancing considerations in its favor. By contrast, bold skepticism can appear to be supported by skeptical arguments like AI. Though the bold skeptic's resolution of our puzzle involves us in systematic falsehood because of its unwavering acceptance of AI's conclusion, it at the same time can seem to make sense of *other* pieces of the puzzle (that we're inclined to say that we don't know that skeptical hypotheses are false and to say that we don't know various ordinary things if we don't know these hypotheses to be false), making the warranted assertability maneuver seem more motivated here than it is in the hands of our imagined crazed philosopher. But, as we saw in the previous section, this appearance is deceptive. Bold skepticism, by itself, does not explain the plausibility of AI's premises. To help the skeptic solve the puzzle, I've had to ascribe to him an analogue of our new solution.⁴⁶ But once we see that the skeptical puzzle can be solved just as well without the bold skeptic's systematic falsehood, we're left with no reason for paying that high price for a solution.⁴⁷ Indeed, since the bold skeptical solution and our

⁴⁶Of course, skeptics are free to refuse this help and propose other solutions. Like practically any claim to have provided the best explanation of something, my claim here is hostage to the possible future development of a better explanation coming along.

⁴⁷Well, little reason. In his 1984, as part of his case for his relativist conclusion that there's no fact of the matter as to whether contextualism or skeptical invariantism is correct, Unger tries to balance this relative disadvantage of skeptical invariantism against contextualism's relative disadvantage that it does not make the truth conditions of knowledge attributions appropriately independent from the current intents and interests of those who happen to be speaking on a given occasion (37). In part 3 of my 1992, I argue that contextualism can handle the most serious consequences one might suspect would follow from this lack of independence. Whatever independence concerns might remain with contextualism seem quite swamped by the cost of the bold skeptic's solution, which, as I've here argued, is quite high indeed.

In his review of Unger 1984, Brueckner, relating the advantages of invariantism, writes, "In particular, speakers' intuitions concerning the correct use of 'know' seem to conform to the closure principle for knowledge asserted by the invariantist yet denied by the contextualist" (1986, 512). If invariantism, but not contextualism, upheld closure, I would take this to be a very important advantage for invariantism—perhaps even weighty enough to make the contest between the two theories interesting. But, as I've argued, contextualism need not, and, properly developed, does not,

new contextualist solution under consideration closely parallel each other, there's not much difference in how they solve the puzzle. That the bold skeptical resolution involves us in systematic falsehood is one of the few differences to be found here, and it's a weighty consideration against that resolution. And, with there being little room for weighty compensating advantages for this resolution over the contextualist's (given how similar they are in other respects), this consideration proves decisive. So, as with the crazed philosopher's theory of 'physician', the bold skeptic's resolution of AI should be rejected because it involves us in systematic and widespread falsehood in our use of a common term of our language.

17. Begging the Question Against the Skeptic?

If skeptics are allowed to play King of the Mountain—they start off on top (never mind how they got there) and it's the anti-skeptics' job to knock them off—displacing them can be a very difficult task. How difficult depends on several factors, one of which is what premises the anti-skeptic is allowed to appeal to in an argument designed to dethrone the skeptic. If the skeptic won't allow any premises to be available, then, as Thomas Reid noted, "It would be impossible by argument to beat him out of this stronghold; and he must even be left to enjoy his scepticism" (1895, 447).⁴⁸ If, to make the game a bit more interesting, a slim range of claims is allowed to pass inspection and be available for use in the anti-skeptical campaign, then (as Reid again recognized) it's often difficult to say what, if anything, of importance would follow from the fact that the skeptic can or cannot be knocked from his perch by arguments from premises of that particular type.

I have little interest in playing King of the Mountain. But skeptical arguments like AI threaten to show that the skeptic needn't just play this game, but can *gain* the top of the mountain—that starting from our own beliefs and intuitions, he can give us better reasons for accepting his skepticism than we have for rejecting it. I've here argued that the bold skeptic cannot win *this* battle—that of providing the best resolution of our puzzling conflict of intui-

take an implausible stand on the issue of closure. (See section 10 and especially note 33, above.)

⁴⁸I discuss this in section II.B of my 1989.

tions. Although AI's premises are initially plausible, the best resolution for the conflict of intuitions generated by AI is not that of the bold skeptic.

Along the way, I've been assuming certain things that we believe but that the skeptic claims we can't know, thereby perhaps raising the concern that I'm begging the question against the skeptic. For instance, in claiming that my belief that I have hands is sensitive, I betray my conviction that I'm not a BIV, either in the actual world or in any nearby worlds. Indeed, I'm ready to admit to the skeptic that if I am a BIV, then I don't know I have hands, according to any standards for knowledge. But, of course, as I firmly believe, I'm not a BIV.

Is it legitimate for me to use this conviction in a debate against the skeptic? Not if we're playing King of the Mountain. But if the skeptic is marshalling deeply felt intuitions of ours in an attempt to give us good reasons for accepting his skepticism, it's legitimate to point out that other of our beliefs militate against his position, and ask why we should give credence to just those that favor him. And if we can further show that those beliefs that seem to favor his solution can be accommodated in our solution better than he can accommodate those of our beliefs that are hostile to him, the best conclusion we can draw is that we're *not* ordinarily mistaken when we claim or ascribe knowledge, despite the bold skeptic's attempt to show that we are. Instead, the main insights to be drawn from a study of AI involve the context-sensitivity of attributions of knowledge, and the role that the Rule of Sensitivity plays in changing the epistemic standards that govern these attributions.⁴⁹

Rice University

⁴⁹This paper is dedicated to the memory of Ken Konyndyk, my first philosophy teacher. Thanks to Anthony Brueckner, John Carroll, Graeme Forbes, Richard Grandy, Mark Heller, Edward Stein, Holly Thomas, an anonymous reader for the *Philosophical Review*, and to audiences at Rice University and at the University of Virginia, Charlottesville for helpful comments on earlier drafts of this paper. Special thanks are due to Peter Unger for his important writings on this topic, from which I've learned much; for three years of almost daily philosophical discussions, many of which were on the topic of knowledge and skepticism and almost all of which were enjoyable and enlightening; and for his many comments on various drafts of this paper and its ancestors.

References

- Brueckner, Anthony L. 1985. "Skepticism and Epistemic Closure." *Philosophical Topics* 13:89–117.
- . 1986. Review of Unger, *Philosophical Relativity*, *Journal of Philosophy* 83:509–17.
- Cohen, Stewart. 1987. "Knowledge, Context, and Social Standards." *Synthese* 73:3–26.
- . 1988 "How to be a Fallibilist." *Philosophical Perspectives* 2:91–123.
- DeRose, Keith. 1989. "Reid's Anti-Sensationalism and His Realism." *Philosophical Review* 98:313–48.
- . 1990. "Knowledge, Epistemic Possibility, and Scepticism." Ph.D. diss., University of California, Los Angeles.
- . 1991. "Epistemic Possibilities." *Philosophical Review* 100:581–605.
- . 1992. "Contextualism and Knowledge Attributions." *Philosophy and Phenomenological Research* 52:913–29.
- Dretske, Fred. 1970. "Epistemic Operators." *Journal of Philosophy* 67:1007–23.
- . 1971. "Conclusive Reasons." *Australasian Journal of Philosophy* 49: 1–22.
- . 1981a. "The Pragmatic Dimension of Knowledge." *Philosophical Studies* 40: 363–378.
- . 1981b. *Knowledge and the Flow of Information*. Cambridge: MIT Press, Bradford Books.
- Forbes, Graeme. 1984. "Nozick on Scepticism." *Philosophical Quarterly* 34: 43–52.
- Goldman, Alvin I. 1976. "Discrimination and Perceptual Knowledge." *Journal of Philosophy* 73:771–91.
- Lewis, David. 1979. "Scorekeeping in a Language Game." *Journal of Philosophical Logic* 8:339–59.
- Luper-Foy, Stephen, ed. 1987. *The Possibility of Knowledge: Nozick and His Critics*. Totowa, N.J.: Rowman & Littlefield.
- Nozick, Robert. 1981. *Philosophical Explanations*, 167–228. Cambridge: Harvard University Press.
- Reid, Thomas. 1895. *The Works of Thomas Reid*, 8th ed., ed. William Hamilton. Edinburgh: James Thin.
- Stine, Gail C. 1976. "Skepticism, Relevant Alternatives, and Deductive Closure." *Philosophical Studies* 29:249–61.
- Stroud, Barry. 1984. *The Significance of Philosophical Scepticism*. Oxford: Oxford University Press.
- Unger, Peter. 1975. *Ignorance: A Case for Scepticism*. Oxford: Oxford University Press.

- . 1984. *Philosophical Relativity*. Minneapolis: University of Minnesota Press.
- . 1986. "The Cone Model of Knowledge." *Philosophical Topics* 14: 125–78.