XIII*—CONTEXTUALIST SOLUTIONS TO SCEPTICISM

by Stephen Schiffer

I

Sceptical Paradoxes. Classical sceptical arguments present us with classical sceptical paradoxes, and the problem of scepticism is the problem of solving those paradoxes. For instance, consider the following sceptical argument that will serve as our exemplar:

[SA] I don't know that I'm not a BIV (i.e., a bodiless brain in a vat who has been caused to have just those sensory experiences I've had).

If I don't know that I'm not a BIV, then I don't know that I have hands.

I don't know that I have hands.[1]

This argument presents a paradox because it tempts us to say three things that are mutually inconsistent: its first premise is true; its second premise is true; and its conclusion is false.

It's natural to suppose that a fully satisfactory solution to this paradox would do two things: first, tell us which of the mutually inconsistent propositions is really false, and second, explain to us why this false proposition *appeared* to be true.[2]

1. This formulation is borrowed from Keith DeRose, 'Solving the Skeptical Problem', *The Philosophical Review*, 104 (January 1995): 1–52. All DeRose citations are to this work.

2. This second constraint on an 'evidently fully satisfying solution' (the reason for my coyness will be revealed later) is well appreciated by DeRose:

> [W]e should hope for a better treatment of the argument than simply choosing which of the three individually plausible propositions—the two premises and the negation of the conclusion—seems least certain and rejecting it on the grounds that the other two are true. In seeking a solution to this puzzle, we should seek an explanation of how we fell into this skeptical trap in the first place, and not settle for making a simple choice among three distasteful ways out of the trap. We must explain how two premises that together yield a conclusion we find so incredible can themselves seem so plausible to us. Only with such an explanation in place can we proceed with confidence and with understanding to free ourselves from the trap. (p. 3)

*Meeting of the Aristotelian Society, held in the Senior Common Room, Birkbeck College, London, on Monday, 20th May, 1996 at 8.15 p.m.

I begin this paper by discussing an interesting resolution of the sceptical paradox which is based on so-called 'contextualist' accounts of the semantics of knowledge sentences. I'll argue that while the contextualist fails to solve the paradox, an appreciation of where she is going wrong may suggest a better resolution. Since more than one kind of semantics for knowledge sentences may be called contextualist, I'll refer to those contextualists I intend to single out as *C*ontextualists.

II

Contextualist Semantics for Knowledge Sentences. Semantic accounts of knowledge sentences don't all share the same primary motivation. There is, for example, the project of trying to shore up the justified-true-belief account of knowledge so that it will withstand Gettier counterexamples. And there are Contextualist accounts of knowledge claims, our present concern, which have been motivated by the desire to answer the sceptic in a way that gives the sceptic her due while, at the same time, protecting the commonsense view of ordinary knowledge claims.[3] The details of any particular Contextualist account won't matter for my purposes, but when I need an illustration, I'll rely on Keith DeRose's already-cited article. I'll first give the essence of a Contextualist semantics for knowledge sentences, and then, in the next section, I'll bring that to bear on the paradox [SA] presents.

According to Contextualism, a sentence of the form '*S* knows that *p*' has no complete proposition as its meaning; different utterances of the sentence, in different contexts of utterance, can express different propositions. Of course, the different propositions expressed by utterances of '*S* knows that *p*' have a good deal in common: they all entail that *p* is true, that *S* believes that *p*, and that *S* satisfies contextually-relevant standards for knowledge. What varies are the standards for knowledge, standards that specify how strong an epistemic position *S* must be in with respect to *p* in order to count as knowing that *p*. Here we may simplify harmlessly. What's expressed in certain contexts is that *S* knows that *p* relative

3. In addition to DeRose, op. cit., see also: David Lewis, 'Scorekeeping in a Language Game', *Journal of Philosophical Logic* 8 (1979): 339–59, and 'Elusive Knowledge,' forthcoming; Stewart Cohen, 'Knowledge, Context, and Social Standards', *Synthese* 73 (1987): 3–26, and 'How to Be a Fallibist', *Philosophical Perspectives* 2 (1988): 91–123; Peter Unger, 'The Cone Model of Knowledge', *Philosophical Topics* 14 (1986): 125–78.

to standard Easy, and what's expressed in certain other contexts is that *S* knows that *p* relative to standard Tough. Standard Tough requires a very strong epistemic position, one that's not met either by one's epistemic position with respect to the proposition that one's not a BIV or by one's epistemic position with respect to the proposition that one has hands. Standard Easy requires a relatively weaker epistemic position, one that can readily be met by both of those epistemic positions. This is a simplification in two respects. First, a reasonable Contextualism won't have just two standards. One's epistemic position can have varying degrees of strength, and different contexts of utterance may invoke differences of degree. Second, the propositional form *S knows that p relative to standard N* is merely intended as a convenient mnemonic; a finished Contextualist theory would doubtless settle on a more structurally perspicuous propositional form.

The notion of the strength of an epistemic position is intuitive but vague. Presently I'll mention DeRose's instructive gloss of this notion, but its vagueness won't hamper me, since I'm happy to assume that, however the Contextualist glosses the standards Tough and Easy, her claims about what epistemic states do or do not satisfy those standards are correct. We do, however, need to note that the Contextualist holds that one's epistemic position with respect to the proposition that one isn't a BIV is at least as strong as one's epistemic position with respect to the proposition that one has hands. 'I am in no better a position to know that *O* [e.g., the *o*rdinary proposition that I have hands] than I am in to know that not-*H* [e.g., the denial of the sceptical *h*ypothesis that I'm a BIV]. This comparative fact is revealed...by the highly plausible conditional...: If I don't know that not-*H*, then I don't know that *O*. Closely tied to that comparative fact...is the related and intuitively compelling realization that it would be no wiser to bet one's immortal soul on *O*'s being true than to bet it on not-*H*'s being true' (DeRose, p. 32).

This brings us to a further defining feature of Contextualism. Some philosophers would reply to [SA] by denying its second premise.[4] These philosophers deny that knowledge is closed under

4. See, e.g., Fred Dretske, 'Epistemic Operators,' *The Journal of Philosophy*, 67 (1970): 1007–23, and Robert Nozick, *Philosophical Investigations*, Cambridge, Mass.: Harvard University Press (1981).

known implication; that is, they deny that if one knows both that p and that p implies q, then one knows that q. The Contextualist quite rightly wants no part of this.[5] But in accepting the closure principle, the Contextualist has to adapt it to his own account. When knowledge is seen as relative to variable standards of knowledge, the epistemic closure principle must take this relativity into account. In the sorts of cases of relevance to us, when one knows that p implies q (e.g., that I have hands implies that I'm not a BIV), then one knows this in a way that conforms to the strictest standards of knowledge. Let 'knows*' be that sense of knowledge. Then a presently convenient way to state the idea that knowledge is closed under known implication would be along the lines of the following relativized closure principle:

> If x knows that p relative to standard N and x knows* that if p, q, then x knows that q relative to standard N.

But what are we then to make of the *sentence* 'If x knows that p, then x knows that q' when x knows* that if p, q? Well, that has to depend on what standards are implicitly appealed to in the antecedent and consequent. There are the following four propositions that could be meant:

> If x knows that p relative to Easy, then x knows that q relative to Easy.
>
> If x knows that p relative to Tough, then x knows that q relative to Tough.
>
> If x knows that p relative to Tough, then x knows that q relative to Easy.
>
> If x knows that p relative to Easy, then x knows that q relative to Tough.

The first two propositions have their truth secured by the closure principle together with the supposition that x knows* that p implies q. The third proposition has its truth secured by the closure principle together with the supposition and the fact that satisfaction of Tough entails satisfaction of Easy. But the fourth proposition may well be false.[6]

5. See DeRose's very nice 'abominable-conjunction' objection to denying the closure principle, pp. 27–29.

A final feature of the Contextualist semantics for knowledge sentences concerns the way in which the context of utterance helps to determine the proposition expressed by the utterance of a knowledge sentence. Specifically, Contextualists want to make two crucial claims, one regarding quotidian contexts, the other regarding contexts in which sceptical hypotheses are at issue. The former claim is that an utterance of '*S* knows that *O*' (where *O* is an 'ordinary' proposition like *I have hands*) in a quotidian context will typically express the true proposition that *S* knows that *O* relative to Easy. The latter claim is that Tough is the prevailing standard in those contexts where sceptical hypotheses are at issue. This claim has two important corollaries. The first is that any claim to know, or not to know, that one's not a BIV automatically creates a context in which Tough is the operative standard, so that the *only* proposition expressible by an utterance of 'I don't know that I'm not a BIV' is the true proposition that one doesn't know that one's not a BIV relative to Tough. The second corollary is that when a sceptical hypothesis has been made the topic, then that generates an ongoing conversational context in which Tough prevails, so that a subsequent utterance, in that context, of 'I know that I have hands' will express the false proposition that one knows, relative to Tough, that one has hands, and a subsequent utterance, in that context, of 'I don't know that I have hands' will express the true proposition that one doesn't know, relative to Tough, that one has hands.

6. Here I'm implicitly correcting DeRose, whose own position on the closure principle seems confused given his Contextualist semantics of knowledge ascriptions. His position on the *sentence*

> If *x* knows that *p*, then *x* knows *q*

when *x* knows* that *p* implies *q* is that it can never be false owing to the fact that *x*'s epistemic position with respect to *q* is at least as strong as *x*'s epistemic position with respect to *p*. He says that such conditionals are true 'regardless of how high or low the standards of knowledge are set' (35), and he gives an analogy:

> Just as the comparative fact that Wilt is at least as tall as Mugsy has the result that the conditional *If Mugsy is tall, then Wilt is tall* will be true regardless of how high or low the standards for tallness are set, so the comparative fact that *x* is in at least as strong an epistemic position with respect to *q* as *x* is in with respect to *p* will result in *If x knows that p, then x knows that q* being true regardless of how high or low the standards for knowledge are set. (35; I've changed DeRose's wording slightly to make it conform to my example.)

But it's simply false that if Wilt is at least as tall as Mugsy, then the conditional *If Mugsy is tall, then Wilt is tall* must be true. A sentence of the form '*a* is tall' expresses a proposition of the form *a is tall for an X*, and consequently nothing precludes an utterance of the conditional from expressing the *false* proposition that if Mugsy is tall for a midget, then Wilt is tall for an NBA center.

Contextualists differ on what they take to be the *mechanisms* by which contexts help to determine standards of knowledge, but I think the best account is DeRose's, and we may usefully flesh out the foregoing sketch of Contextualist semantics by briefly considering his account of the standard-inducing mechanism and how this relates to the notion of epistemic strength.

Inspired by what he takes to be the truth in Nozick's (op. cit.) subjunctive-conditionals account of knowledge, DeRose appeals to the notion of the *sensitivity/insensitivity* of a belief state. Subject to certain qualifications that I'll ignore,[7] we may say that:

> *x*'s belief that *p* is *insensitive* iff *x* would have *x*'s belief that *p* even if *p* were false.

Thus, my belief that I'm not a BIV is insensitive, since I would believe that I'm not a BIV even if I were a BIV; that is to say, in the nearest possible worlds in which I am a BIV, I believe that I'm not a BIV. But my belief that I have hands is sensitive, since if I didn't have hands, I'd believe that I didn't; that is to say, in the nearest possible worlds in which I don't have hands (e.g., possible worlds in which I lost my hands in an accident), I believe that I don't have hands. Notice that, even though my belief that I have hands is sensitive and my belief that I'm not a BIV isn't, it still remains true that my epistemic position with respect to the proposition that I'm not a BIV is at least as strong as my epistemic position with respect to the proposition that I have hands. DeRose roughly glosses this last point in the following way (pp. 33–35). I'm in a pretty strong epistemic position with respect to the proposition that I have hands because my belief as to whether or not I have hands tracks the truth pretty well over most nearby, and not so nearby, possible worlds. My belief that I have hands is sensitive, because the nearest possible world in which I don't have hands isn't all that far away; it's in the sphere of my reliable tracking. Of course, in the possible world in which I'm a BIV, then my belief that I have hands doesn't track the truth, but that's a *very distant* possible world, and my failure to track the truth there doesn't preclude me from being in a pretty strong epistemic position with respect to the proposition that I have hands. Now, in all but very distant possible worlds, I'm not a BIV, so my epistemic

7. See the discussion of sensitivity and Nozick's account that runs throughout DeRose's article.

position with respect to the proposition that I'm not a BIV is pretty strong simply by virtue of my believing that proposition in all those nearby and not-so-nearby possible worlds in which I'm not a BIV. Thus, we can reconcile the fact that my belief that I'm not a BIV is insensitive with the fact that I'm in a pretty strong epistemic position with respect to the proposition that I'm not a BIV: I continue to believe that I'm not a BIV in the nearest possible world in which I am a BIV, but that's still a *very distant* possible world.

Equipped with the notion of sensitivity, DeRose proposes a standard-inducing mechanism he calls the *Rule of Sensitivity*:

> When it is asserted that some subject *S* knows (or does not know) some proposition *P*, the standards for knowledge (the standards for how good an epistemic position one must be in to count as knowing) tend to be raised, if need be, to such a level as to require *S*'s belief in that particular *P* to be sensitive for it to count as knowledge. (p. 35)

A consequence is that:

> Where the *P* involved is to the effect that a skeptical hypothesis does not obtain, then this rule dictates that the standards will be raised to a quite high level, for...one must be in a stronger epistemic position with respect to a proposition stating that a skeptical hypothesis is false—relative to other, more ordinary, propositions—before a belief in such a proposition can be sensitive. (p. 35)

We're now in a position to consider the Contextualist's solution to scepticism.

III

The Contextualist Solution to Scepticism. The Contextualist's response to the sceptical argument [SA] is determined by her semantics for knowledge sentences. She holds that a sentence of the form '*S* knows/doesn't know that *p*' doesn't express the proposition *that S knows/doesn't know that p*, for there is no such proposition. The proposition expressed must be of the form *S knows/doesn't know that p relative to standard N*, and we're provisionally pretending that Tough and Easy are the only relevant standards. At the same time, the Contextualist holds that the mechanism by which context determines the proposition expressed by a knowledge sentence (e.g., DeRose's Rule of Sensitivity)

secures that when *p* is a sceptical hypothesis or its denial, then the standard must be Tough, and thus the *only* proposition expressible by [SA]'s first premise is the *true* proposition that I don't know that I'm not a BIV relative to Tough.[8] Now, the Contextualist further holds, the sceptic's mere utterance of the first premise, 'I don't know that I'm not a BIV', induces Tough *for the 'context' that includes the whole of her argument*. That is, by the mere assertion of her first premise, the sceptic raises the standard for knowledge to Tough for *the entire argument*. Consequently, the argument *really expressed* by [SA] is the *sound* argument:

> I don't know that I'm not a BIV relative to Tough.
>
> If I don't know that I'm not a BIV relative to Tough, then I don't know that I have hands relative to Tough.
>
> ---
>
> I don't know that I have hands relative to Tough.

The Contextualist is pleased with this verdict because it ostensibly accounts for the force of the sceptic's argument while allowing that many of the knowledge claims we make in quotidian contexts, where scepticism isn't at issue and Easy prevails, are true.

We apparently have a neat (perhaps too neat) dissolution of the paradox [SA] *seemed* to provide. A paradox is a set of mutually inconsistent propositions each of which enjoys some plausibility when considered on its own. The paradox that [SA] *seemed* to generate was, we were naively tempted to say, the set consisting of these three propositions: that I don't know that I'm not a BIV; that if I don't know that I'm not a BIV, then I don't know that I have hands; that I know that I have hands. What the Contextualist must now say is that we don't really have a paradox, a set of mutually inconsistent *propositions* each of which enjoys some plausibility considered on its own. For the set of mutually inconsistent *propositions* we now see that we have contains *these* three propositions: that I don't know that I'm not a BIV relative to Tough; that if I don't know that I'm not a BIV relative to Tough, then I don't know that I have hands relative to Tough; that I know that I have hands relative to Tough. Yet this fails to be a *paradox*,

8. It's definitive of Contextualism that one knows neither that not-*H* nor that *O* relative to Tough, but each Contextualist owes her own account of Tough to show why this is so. For DeRose, Tough is induced by the Rule of Sensitivity, and thus requires an epistemic position so strong that it continues to track the truth even in those very distant worlds where the sceptical hypothesis *H* is true.

for the third proposition has no plausibility when viewed on its own; it's plainly false.

Naturally, the Contextualist realizes that he isn't quite finished. He must explain why [SA] *seemed* to present a paradox. If the argument [SA] *really* expresses is plainly sound, then why do we instinctively feel that [SA] expresses an argument that's plainly unsound? Why, that is, are we loathe to accept that the sceptic's conclusion is true? To this the Contextualist has a simple answer: we instinctively know that the conclusion-asserting sentence of [SA] would express a false proposition in a quotidian context in which sceptical hypotheses weren't at issue, and we mistakenly suppose that it's asserting the same false proposition in [SA]. In other words, [SA] strikes us as presenting a profound paradox merely because we're ignorant of what it's really saying, and this because we don't appreciate the indexical nature of knowledge sentences.

IV

A Problem for Contextualism. The Contextualist's dissolution of the sceptical paradox seemingly generated by [SA] has two parts: first, her claims about the semantics of knowledge sentences, and second, a certain *error theory*—to wit, the claim that people uttering certain knowledge sentences in certain contexts systematically confound the propositions their utterances express with the propositions they would express by uttering those sentences in certain other contexts. Both the semantics and the error theory are needed for the Contextualist's response to [SA]. The semantics is needed to locate the false proposition in the set of mutually inconsistent propositions that we get by combining [SA]'s premises with the denial of its conclusion, and the error theory is needed to explain why the sentence expressing the false proposition—the sentence 'I know that I have hands'—deceptively appears to be stating a true proposition. The trouble is that the semantics is refuted by the error theory.

The semantic claims about knowledge sentences I've attributed to the Contextualist don't add up to a complete semantic account of those sentences. To what sort of completing semantics might the Contextualist appeal? The important thing is to capture the needed 'indexical' nature of knowledge claims, the fact that different utterances of a knowledge sentence with no apparently indexical terms

('Hillary Rodham Clinton knows that Bill Clinton was a Rhodes Scholar') can express different propositions. I can think of three possible accounts, but none of them will help the Contextualist.

Since a knowledge sentence is supposed to express different propositions in different contexts even if it contains no apparently indexical terms, one naturally thinks of a 'hidden-indexical' theory of knowledge sentences. This would be to assimilate knowledge sentences to sentences like 'It's raining' or 'He's short'. Such sentences express propositions that contain what John Perry has called *unarticulated constituents*: propositional constituents that aren't the semantic values of any terms in the uttered sentence.[9] Thus, an utterance of 'It's raining' might express the proposition that it's raining *in London*, and an utterance of 'He's short' might express the proposition that so-and-so is short *for an NBA center.* This would be the most likely account if we could really take seriously the idea that sentences of the form '*S* knows that *p*' express propositions of the form *S knows that p relative to standard N.*

What's hard to see is how the hidden-indexical proposal can sustain the idea that fluent speakers systematically confound their contexts, so that even when they're in a context in which Tough is the induced standard occurring in the *false* proposition *they* have just asserted, they mistakenly think they've just asserted a *true* proposition, a proposition that evidently contains the standard Easy that would be induced by an utterance of the problematic sentence in a quite different context. It's as though a fluent, sane, and alert speaker, who knows where she is, were actually to assert the proposition that it's raining in London when she mistakenly thinks she's asserting the proposition that it's raining in Oxford. Actually, the situation is even much more problematic. For the speaker would not only have to be confounding the proposition she's saying; she'd also have to be totally ignorant of the sort of thing she's saying. One who implicitly says that it's raining in London in uttering 'It's raining' knows full well what proposition she's asserting; if articulate, she can tell you that what she meant and was implicitly stating was that it was raining in London. But no ordinary person who utters 'I know that *p*', however articulate, would dream of telling you that what he meant and was implicitly stating was that

9. 'Thought without Representation,' *Proceedings of the Aristotelian Society, Supplementary Volume*, 60 (1986): 137–51. See also my 'Belief Ascription,' *The Journal of Philosophy*, 89 (1992): 499–521.

he knew that *p* relative to such-and-such standard. If, for example, this ordinary guy says 'I know that Placido Domingo is scheduled to sing at the Met this season' and you ask him what exactly he said, he'll tell you that what exactly he said, and meant, was that he knew that Placido Domingo was scheduled to sing at the Met this season.

A second way of accommodating the alleged context-sensitivity of knowledge sentences would be to claim that the verb 'to know' was itself indexical. Relative to the pretence that the Contextualist recognizes only two knowledge standards, Easy and Tough, this would be tantamount to holding that an utterance of 'I know that I have hands' may express, according to the context of utterance, either the proposition that I know_e that I have hands or the proposition that I know_t that I have hands, where the knowing_e relation encapsulates the standard Easy and the knowing_t relation encapsulates the standard Tough.[10] But I can't see that it's any improvement to say that those uttering knowledge sentences are both referring, unbeknown to themselves, to different knowledge relations and confounding the knowledge relations to which they're unknowingly referring.

The hidden-indexical and indexical-verb proposals are implausible even apart from the further error aspect that increases their implausibility. A final proposal is not at all implausible on its own, though it can't be combined with the error theory to which the Contextualist is forced to appeal. The verb 'to know', like virtually every expression, is vague, and there is a certain context variability inherent in vagueness. The penumbras of vague terms can dilate or constrict according to conversational purposes. For example, in a conversation about good places to run, a runner might say 'What I especially like about Hyde Park is that, unlike Central Park, it's flat' and count as speaking truly, whereas in a conversation among engineers about where a certain flying device might land in an emergency, an engineer might count as speaking truly when he says of Hyde Park 'It's not flat'. It's clear that knowledge sentences are subject to this sort of vagueness-related variability. In certain conversational contexts you count as knowing that your spouse is faithful; in others you can't really be said to *know*. The reason this

10. Propositional-attitude verbs are indexical on the account developed in Mark Richard, *Propositional Attitudes* (Cambridge: Cambridge University Press, 1989); but for Richard the variable component has to do with modes of presentation rather than epistemic standards.

sort of variability is of no use to the Contextualist is that speakers are perfectly aware of when it's going on. If you claim to know where your car is, and someone challenges you, 'But how can you be sure it wasn't towed or stolen?', you'll merely get impatient at the questioner's obtuseness: it ought to have been mutual knowledge between you that you were speaking casually. The Contextualist who tries to appeal to the context-variability of vagueness must say the following. In a context where scepticism is at issue, the penumbra of 'know' shrinks dramatically to such an extent that one who says 'I know that I have hands' is actually making a false assertion. But the speaker, either because she is unaware of how the vagaries of vagueness affect speech or because she confounds her context, mistakenly thinks that more generous precisifications are in play which count her as speaking truly. This is not a semantic story to be taken seriously.

I conclude that, as far as I can see, there is no plausible semantic theory that will resolve sceptical paradoxes in the way the Contextualist requires. If the proposed semantics were correct, then the extreme error theory would be needed to explain why we appear to have a paradox in the first place. But that error theory has no plausibility: speakers would know what they were saying if knowledge sentences were indexical in the way the Contextualist requires.

V

An Unhappy-Face Resolution. A paradox is a set of mutually inconsistent propositions each of which enjoys some plausibility when considered apart from the others. A *happy-face* solution to a paradox simply identifies one of the propositions as the false member of the set, the odd-guy-out. Typically, this odd-guy-out is identified on the basis of a meaning analysis of the crucial concept generating the paradox. It's claimed that once we get clear about the meaning of that concept, we'll see that such-and-such particular proposition is false. Compatibilism, for example, is a happy-face solution to the paradox of free will. There the paradox consists of the following set of mutually inconsistent propositions: (i) we have free will; (ii) everything we do was caused by events that occurred even before we were born; (iii) (i) and (ii) are incompatible. The compatibilist rejects (iii). He says that once we appreciate that to say that an action was done of one's free will is

just to say that it was caused by one's desires and beliefs in a certain way, then we'll see that it's plainly consistent with one's having free will that one's beliefs and desires have causes that antedate one's birth. The typical trouble with happy-face solutions, a trouble well illustrated by the compatibilist's solution to the problem of free will, is that it leaves one wondering, 'If that's the solution, then what the hell was the *problem*?' The happy-face solution makes it a mystery why one was ever deceived by the false proposition in the first place. If '*x* acted freely' simply *means*, say, that *x* would have acted differently if *x* had decided to, then why would anyone ever have thought that determinism threatened free will? Are fluent and intelligent speakers of English confused about what their own words mean in their own mouths?

The Contextualist is in effect attempting a happy-face solution to the sceptical paradoxes. But she's mindful of the trap such a solution can fall into. Whence her error theory: we mistakenly think that 'I don't know that I have hands' expresses a false proposition when the sceptic utters it to assert the conclusion of her argument, and we make this mistake because we're distracted by our knowledge that an utterance of the sentence would express a falsehood in a quotidian context. Yet we've seen that while this error theory is an inevitable corollary of the semantics the Contextualist needs to sustain her happy-face solution, it's a pretty lame account of how, according to her, we came to be bamboozled by our own words.

I think that the argument actually expressed by [SA] presents a genuine paradox; no amount of getting clear about the propositions constituting the premises and conclusion of the argument will dissolve the sense of paradox. Since it's obvious that the argument is valid, a happy-face solution to it would tell us either that the argument is sound or else that one of its two premises is false. In other words, the set of mutually inconsistent propositions we get by combining the premises with the denial of the conclusion really is a paradox, and a happy-face solution to it would have to pick one of the three propositions as the odd-guy-out. I'm inclined to think this can't be done. The paradox has no happy-face solution and requires, therefore, an unhappy-face solution.

We have three mutually inconsistent propositions—that I don't know that I'm not a BIV; that if I don't know that I'm not a BIV, then I don't know that I have hands; that I know that I have hands—

but none is lurking there with its falsity waiting to be discovered, not even by God. The trouble is that our concept of knowledge has a deep-seated incoherence. The criteria we employ to tell us when we know something conflict, and the concept contains no higher criterion whose application can resolve the conflict. To see this, let's briefly survey the apparent licence for asserting each of the three members of our paradox set.

1. *I know that I have hands.* All other things being equal, we perceptually normal people count ourselves as knowing the true propositions we're made to believe by our well-functioning sensory faculties in excellent sensory conditions. The three-year-old child is caused by her tactile, visual, and olfactory sensations to have the true belief that there's an apple in her hand. The child's having this knowledge doesn't depend on her having made any sort of inference from premises about her sense experiences. She's simply built, qua information processor, to believe the output of her perceptual belief-forming mechanisms, and, roughly speaking, we deem these beliefs to constitute knowledge when they're true. When we gain intellectual sophistication, we learn that our perceptual beliefs have potential defeaters: one mustn't be too quick to believe that the thing in the bowl is an apple; this place is known to display wax fruit. But our criteria for ascribing perceptual knowledge are such that a person's perceptual belief counts as knowledge if it's true, his perceptual faculties are well functioning, the belief was formed under excellent perceptual conditions, and there is no special, or particular, reason to believe that any potential defeater obtains.

2. *If I know that I have hands, then I know that I'm not a BIV,* or, equivalently, *if I don't know that I'm not a BIV, then I don't know that I have hands.* Another strand to our concept of knowledge deems one to know what one's deduced with certainty from premises one knows. Some version of the principle that knowledge is closed under known implication is evidently a basic feature of our concept of knowledge. For consider Harold, an intelligent fellow who got an A in Logic 101. He's confronted with an argument of the form $P \rightarrow Q$; P; $\therefore Q$. Harold knows that the conclusion has got to be true if the premises are true, so he can't accept the premises without also accepting the conclusion. What could be better justification for believing Q than deducing it with certainty from premises one knows to be true? So if Harold does

indeed know the premises, then his belief that Q is true and is superbly well justified, and that, by one strand of our concept of knowledge, counts as knowledge.[11]

3. *I don't know that I'm not a BIV.* Some of the criteria implicit in our concept of knowledge tell us that I know that I have hands, and that if I know that I have hands, then I know that I'm not a BIV. After all, since I know that I have hands, I know that I'm not handless, and I know it's true by definition that a BIV is handless. In this way, I should be entitled to assert that I know that I'm not a BIV. Yet, at the same time, there are conflicting criteria implicit in our concept of knowledge according to which I don't know that I'm not a BIV. It might seem that Nozick and, following him, DeRose have put their finger on what the implicit criterion is: my belief that I'm not a BIV is insensitive—I'd have that belief even if I were a BIV. But, as DeRose is aware, sensitivity, either as glossed above or even as further precisified by Nozick, can't fully account for the plausibility of the sceptic's first premise, her claim that she doesn't know that she's not a BIV. Further refinement is needed, and it's not clear how it should go. The need for refinement is illustrated by the sceptical hypothesis that I'm a BIV′, a BIV′ being exactly like a BIV except that it lacks auditory sensations. Now I certainly believe that I'm not a BIV′, and my belief that I'm not a BIV′ is insensitive: if I were a BIV′, I'd still believe that I wasn't. But it's obvious that I *know* that I'm not a BIV′, for I know that I'm having auditory sensations and that I wouldn't if I were a BIV′.

But wait! If, despite the insensitivity of my belief, I can straightforwardly and unproblematically know that I'm not a BIV′ on the basis of knowing that I have auditory sensations, then why can't I simply know, in the same way, that I'm not a BIV on the basis of knowing that I have hands? The answer, I venture, has to do with the fact that being a BIV is a potential defeater for one's belief that one has hands, whereas being a BIV′ isn't a potential defeater for one's belief that one's having auditory sensations. My belief that I'm having auditory sensations is for all relevant purposes rationally indefeasible; but my belief that I have hands can be rationally defeated in various ways, among which would be

11. To be sure, Harold might deduce Q and, finding this questionable, proceed to question the hitherto unquestioned premises. But *if* he knows that p and that if p, q—and he's not psychologically prevented from putting two and two together—then he knows that q.

by my gaining evidence that my sensory experiences were caused in some way other than by my causally interacting with hands that were mine. Now here's where the problem lies. If *H* is a potential defeater of one's belief that *p*, then we take it to be question begging to attempt to refute *H* by appeal to the fact that *p*. What's needed is a basis for rejecting *H* that bears directly on it. The reason one seems not to know one's not a BIV, to the extent that one seems not to know it, is that the directly-bearing reasons seem hard to find. We accept the External World hypothesis as the correct explanation of the fact that we have the sensory experiences we have, but when we're forced to consider the BIV hypothesis as an alternative explanation, we find that there's no basis for preferring the External World hypothesis that adds up to our *knowing* that the BIV hypothesis isn't correct.[12]

The bind that our concept of knowledge puts us in may be looked at in a slightly different light. Our perceptual faculties provide innate belief-forming mechanisms that are noninferential. Certain kinds of sensory hits under certain kinds of conditions cause us to have certain beliefs about our perceptual environment. At a certain stage in our cognitive development, higher-level belief-forming mechanisms also take root; we learn how to reason and to theorize about our own place in the world. It's at this point that our perceptual beliefs become defeasible. One appreciates that for one's perceptual belief to count as knowledge it must have a certain etiology, and that all that one knows 'directly'—viz., what one believes and what sensory experiences one's currently having—is consistent with one's perceptual belief's having an etiology that would render the belief false. In practice, one doesn't count this bare possibility as defeating one's claim to know what one perceptually believes; for one's claim to know to be defeated, one needs *special* reason to believe that the defeating etiology obtains. So far so good. The catch is that, at the same time, we don't count ourselves as *knowing* that the defeating etiology doesn't obtain unless we have evidence for its not obtaining that goes beyond what one normally has when one takes oneself to have perceptual knowledge. This is the source of the incoherence in our concept of knowledge. No wonder we're of two minds about the sceptic's claim not to know

12. Cf. Ümit D. Yalçin, 'Skeptical Arguments from Underdetermination', *Philosophical Studies*, 68 (1992): 1–34.

that she's a BIV, and no wonder deep analysis of our concept of knowledge provides no release for this ambivalence. It's why the sceptical paradox has no happy-face solution.

The idea of an unhappy-face resolution to a paradox is of course not new; Tarski, for example, evidently thought it was what had to be said about our commonsense notion of truth as regards the semantic paradoxes,[13] and this has most recently been spelled out in some detail by Charles Chihara.[14] Now, in the case of our concept of truth, it's not good enough to pronounce our commonsense concept incoherent and be done with it; the further question arises as to how to rehabilitate the notion of truth so that it doesn't lead us into absurdity. The concept of knowledge is not, I dare say, as important as our concept of truth; we could probably live without it. But in recognizing the implicit incoherence of our concept of knowledge, we are faced with an important epistemological question: What *should* our cognitive economy be like? What kind of belief-forming policies ought we to have, and how will they play out with respect to sceptical hypotheses? I can't go into this now, but it's clear that the sceptic will gain no new foothold here.[15]

Department of Philosophy
City University of New York, Graduate Center
33 W42nd Street
New York
NY 10036, USA

13. Alfred Tarski, 'The Concept of Truth in Formalized Languages'. In *Logic, Semantics, Metamathematics*, tr. J. Woodger (Oxford: Oxford University Press, 1956).

14. 'The Semantic Paradoxes: A Diagnostic Investigation,' *The Philosophical Review*, 88 (1979): 590–618. At the close of his paper, Chihara writes:

> As a concluding thought, I should like to note that the inconsistency view of truth advocated in this paper, if correct, has far-reaching consequences for philosophy—especially regarding the widely accepted and generally unstated presupposition of much philosophical work that the fundamental notions that permeate our thinking and that give rise to philosophical perplexities, such as *knowledge*, *evidence*, *rationality*, and *existence*, must be consistent and free from paradox when correctly understood and analyzed. It would seem, as a result of the above analysis of the semantic paradoxes, that this presupposition should be seriously questioned. (p. 618)

My 'unhappy-face' resolution of the sceptical paradox endorses this recommendation. I should also make explicit that my proposal isn't that 'knows that such-and-such', like 'is a round square', means a perfectly coherent property that simply can't be instantiated. It's rather that distinct criteria that purport to govern the use of 'knows' can be made to yield inconsistencies, showing that these ostensibly meaning-determining criteria can't all be correct. Nor does the kind of 'incoherence' in question preclude the notion of knowledge from having useful employment.

15. Thanks to Jonathan Adler, Kent Bach, John Carroll, Russell Dale, Keith DeRose, Hartry Field, Brian Loar, Peter Unger, and Ümit Yalçin for their comments on an earlier draft.